



Flash Memory Summit

Exploiting Managed Language Semantics to Mitigate Wear-out in Persistent Memory

Shoaib Akram

Ghent University, Belgium



Main memory capacity expansion

Charge storage in **DRAM** a scaling limitation

Manufacturing complexity makes **DRAM** pricing volatile





Phase change memory (PCM)

😊 Scalable → More Gb for the same price

Byte addressable like DRAM

Latency closer to DRAM


😢 Low write endurance



Why PCM has low write endurance?

Store information as change in resistance
Crystalline is set & Amorphous is reset

Electric pulses to
program **PCM** cells
wear them out






Flash Memory Summit

Mitigating PCM wear-out

Wear-leveling to spread writes across PCM






Flash Memory Summit

Mitigating PCM wear-out


Wear-leveling to spread writes across PCM





Mitigating PCM wear-out

Wear-leveling to spread writes across PCM



Problem: PCM-Only with wear-leveling wears out in a few months



Flash Memory Summit

Hybrid DRAM-PCM memory

Endurance

Capacity
Persistence

DRAM

PCM

This talk → Use DRAM to limit PCM writes




Flash Memory Summit

OS to limit PCM writes



DRAM



PCM

Page migrations hurt performance and PCM lifetime



Managed runtimes


Platform independence

Abstract hardware/OS

→ Aka *Virtual Machine*

Ease programmer's burden

Garbage collection (**GC**)





GC to limit PCM writes

GC aware of heap semantics
→ Pro-active allocation

GC operates with objects
→ Fine-grained mgmt.

Application



Operating System

Hardware



Flash Memory Summit


Write Distribution in GC heap



70%
of writes



Write Distribution in GC heap





Flash Memory Summit

Write-Rationing Garbage Collection

Limit PCM writes by discovering highly written objects



Kingsguard → dynamic monitoring






Crystal Gazer → prediction




Flash Memory Summit

Kingsguard-Nursery (KG-N)



DRAM




PCM





Flash Memory Summit

Kingsguard-Writers (KG-W)





Metadata optimization

meta

payload

Full-heap **GC**: Mark a bit in meta of all live objects

Meta Opt: Place object meta-data in **DRAM**



KG-W drawbacks

Monitoring overhead

Limited opportunity to predict writes

Fixed DRAM consumption



Flash Memory Summit

Write-Rationing Garbage Collection

Limit PCM writes by discovering highly written objects



Kingsguard → monitoring



Crystal Gazer → prediction



Allocation site as a write predictor

```
a = new Object()  
b = new Object()  
c = new Object()  
d = new Object()
```



Produces highly written objects


Uniform distribution 😞

Skewed distribution 😊



Write distribution by allocation site

Few sites capture majority of writes





Crystal Gazer operation

Application Profiling → Advice Generation → Bytecode Compilation

a = new Object()
...
b = new Object()

a = new Object()
...
b = new_dram Object()



Advice generation

Generate <alloc-site, advice> pairs

advice → DRAM or PCM

input is a write-intensity trace

Two heuristics to classify allocation sites as DRAM



DRAM allocation sites

Frequency: More than a threshold writes

- ✓ Aggressively limits writes
- ✗ 1 Byte and 1024 Byte object treated similarly

Density: More than a threshold *write-density*

- ✓ Optimizes for writes and DRAM capacity



Classification examples

Frequency threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4



Classification examples

Frequency threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4





Classification examples

Frequency threshold = 1

PCM writes = 0/256, DRAM bytes = 5008

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
→ O3	128	4	A() + 10
→ O4	128	4096	B() + 4



Classification examples

Density threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4



Classification examples

Density threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4

→ 32



Classification examples

Density threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4

→ <1



Classification examples

Density threshold = 1

PCM writes = 128/256, DRAM bytes = 12

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4



Object placement in Crystal Gazer


new_dram() → Set a bit in the object header

GC → Inspect the bit on nursery collection to copy object in DRAM or PCM



Flash Memory Summit

Object placement in Crystal Gazer



DRAM



Is marked
highly written? ✓

PCM



Flash Memory Summit

Persistence

Persistent parent → copy child objects to **PCM**

VM startup → Move highly-written to **DRAM**

Write barrier tracks writes & persistent candidates



Evaluation methodology

15 Applications → DaCapo, GraphChi, SpecJBB

Medium-end server platform


Different inputs for production and advice

Jikes RVM



Flash Memory Summit

Emulation platform





Flash Memory Summit


PCM write rates → lifetime

PCM-Only write rate is up to 1.8 GB/s

Safe operation is 200 MB/s for 5-10 year lifetime




PCM write rates






Performance






DRAM capacity





Flash Memory Summit


KG-W versus Crystal Gazer





Crystal Gazer
opens up
Pareto-optimal
trade-offs

KG-W versus Crystal Gazer





Flash Memory Summit

Write-rationing garbage collection

Hybrid memory is inevitable

DRAM

PCM

Each layer can play a role
in wider adoption

Write-rationing GC is
pro-active and fine-grained





More information

PLDI 2018 → Write-rationing garbage collection for hybrid memories

SIGMETRICS 2019 → Crystal Gazer: Profile-driven write-rationing garbage collection for hybrid memories

ISPASS 2019 → Emulating and evaluating hybrid memory for managed languages on NUMA platform