

Emulating and Evaluating Hybrid Memory for Managed Languages on NUMA Hardware

**Shoaib Akram (Ghent), Jennifer B. Sartor (Ghent and VUB),
Kathryn S. McKinley (Google), and Lieven Eeckhout (Ghent)**

Shoaib.Akram@UGent.be



Hybrid DRAM-PCM memory

- 😊 More GB/\$ with Phase Change Memory
- 😢 Higher latency *and low* endurance



DRAM



PCM

Managing DRAM-PCM memory

Mitigate PCM wear-out

Bridge the DRAM-PCM latency gap

Speed
Endurance

Capacity

DRAM

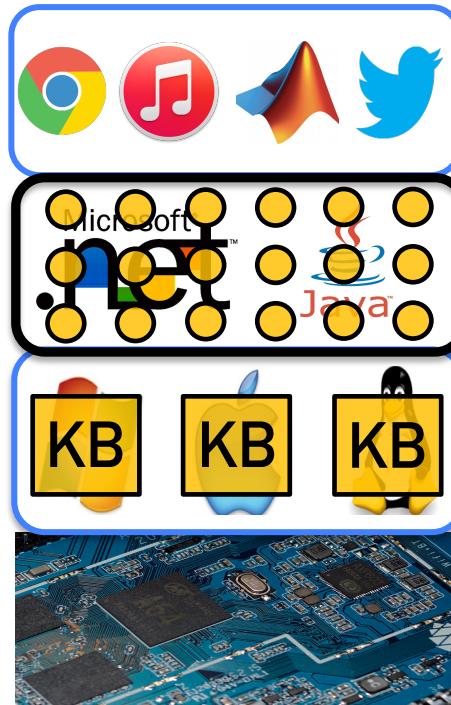
PCM

Managing DRAM-PCM memory



*Write-Rationing
Garbage Collection
for Hybrid Memories*

Operating System
Coarse-grained
pages  KB



Garbage collection
Proactive ☺
Fine-grained
objects ● ● ● ●

GC manages DRAM-PCM hybrid better than OS

Pros/cons of simulating DRAM-PCM

Gain insight

- What triggered the writeback to memory?

Study parameter sensitivity

Slow process

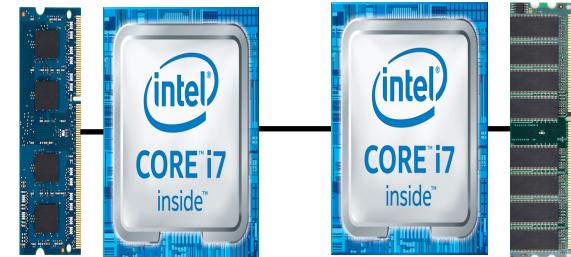
- Page Rank over twitter → hours versus months!

Incomplete model

- Missing OS or proprietary hardware features

Emulation for hybrid memory

Multi-socket NUMA for
emulating DRAM-PCM hybrid
memory



Fast evaluation of emerging workloads
Several co-running BIG graph analytic
applications written in Java

Existing emulation platforms

Focus is to evaluate explicit memory management in **C/C++**

Focus is to model the latency of **PCM**

Contribution: Emulation platform

DRAM-PCM emulation for managed applications



Comparison with Sniper using write-rationing garbage collectors



Contribution: Analysis of PCM writes

PCM writes and write rates

C++ versus Java

Impact of multiprogramming

Classic versus emerging applications

Is PCM practical as main memory?

Outline

Heap management

Kingsguard collectors

Comparison with simulation

Write analysis

Outline

Heap management

Kingsguard collectors

Comparison with simulation

Write analysis

DRAM heap management

Heap Tracker



available



occupied



HEAP_BEGIN

HEAP_END

Heap Organization



DRAM heap management

Heap Tracker

✓ available

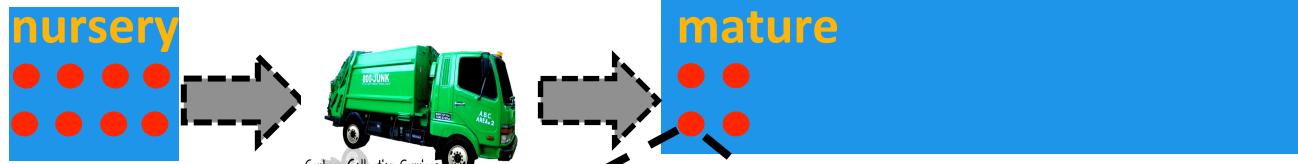
★ occupied



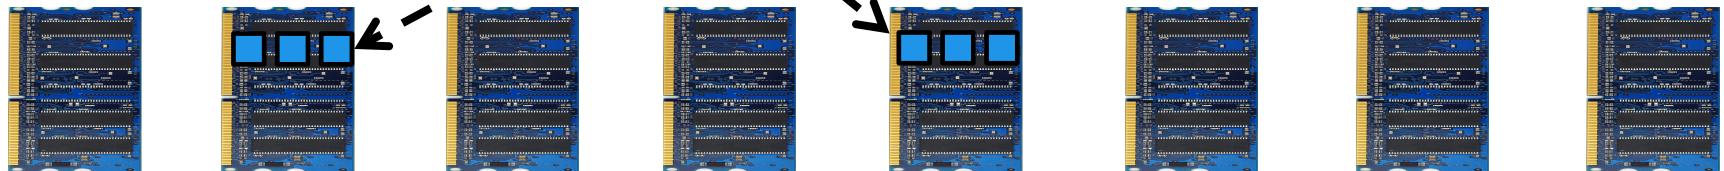
HEAP_BEGIN

HEAP_END

Heap Organization

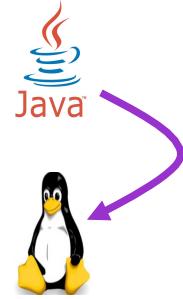


Physical Memory



DRAM-PCM heap management

JVM uses `mbind()` to inform the OS to map a space in **DRAM** or **PCM**



Anything else the JVM should do?

Next: Sanity check with a **DRAM nursery** and **PCM mature**

DRAM-PCM heap management

Heap Tracker

✓ available

★ occupied



HEAP_BEGIN

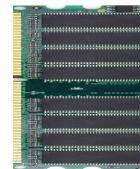
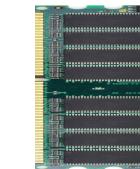
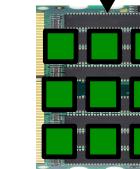
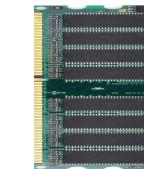
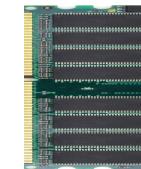
PCM_BEGIN

HEAP_END

Heap Organization



Physical Memory



DRAM-PCM heap management

Heap Tracker

✓ available

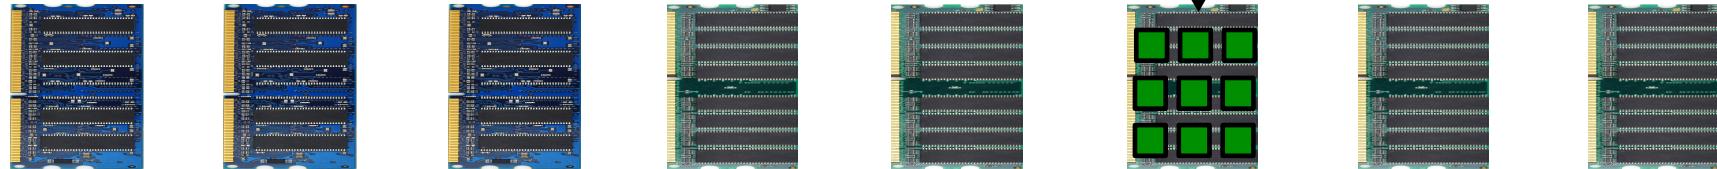
★ occupied



Heap Organization



Physical Memory



DRAM-PCM heap management

Heap Tracker

✓ available

★ occupied



HEAP_BEGIN

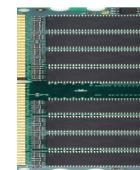
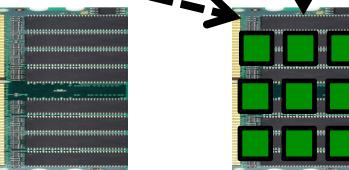
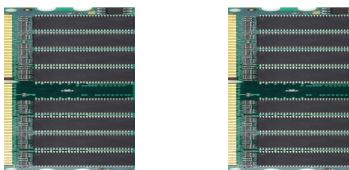
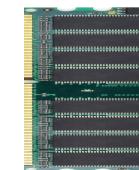
PCM_BEGIN

HEAP_END

Heap Organization



Physical Memory



DRAM-PCM heap management

Options

Map/unmap pages in physical memory whenever space grows/shrinks

Two free lists ✓

DRAM-PCM heap management

Heap Tracker



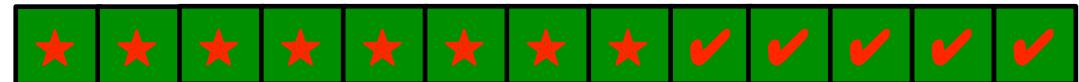
available



occupied



DRAM_BEGIN



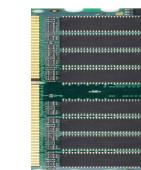
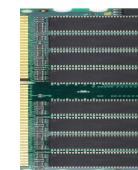
PCM_BEGIN

PCM_END

Heap Organization



Physical Memory



Outline

Heap management

Kingsguard collectors

Comparison with simulation

Write analysis

Kingsguard-Nursery (KG-N)



Write-rationing GC: concentrate writes in DRAM

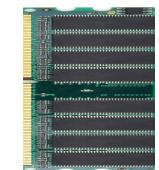
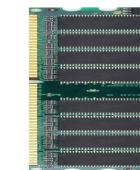
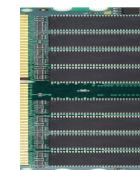
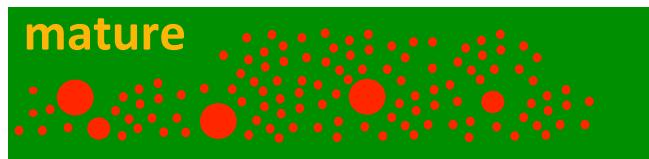
70%

of writes



22%

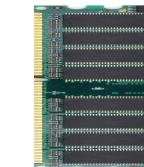
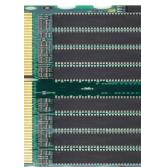
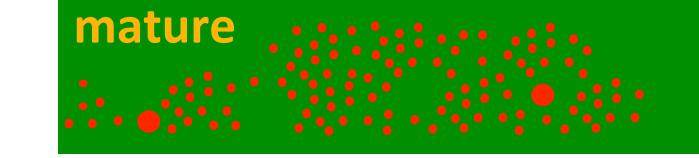
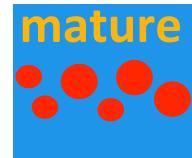
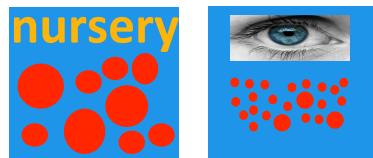
to 2% of objects



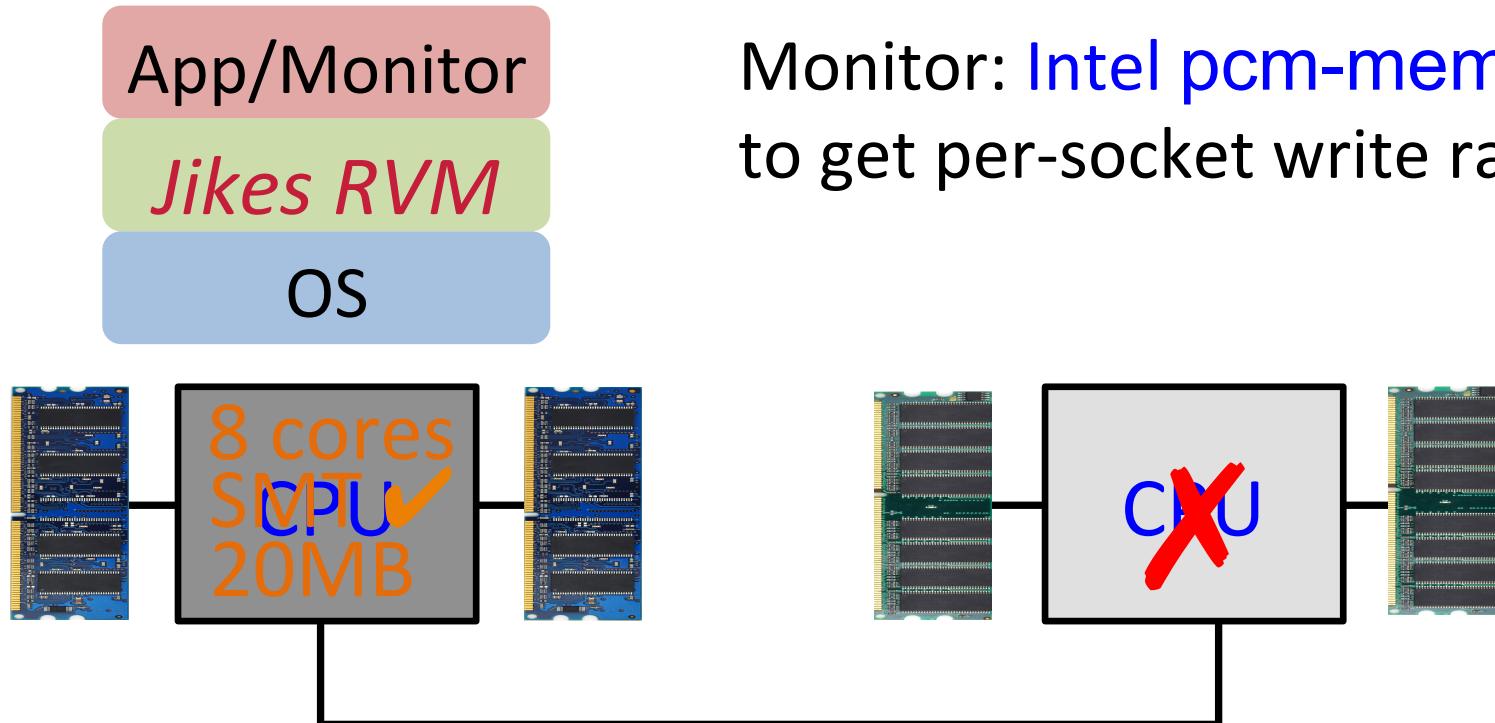
Kingsguard-Writers (KG-W)



KG-W monitors writes in a DRAM observer space
Trades off performance for better endurance



Emulation setup



Emulation versus simulation

PCM write reduction with **KG-N** and **KG-W**
versus **PCM-Only**

Execution time increase with **KG-W** versus **KG-N**

No OS in simulation
Faithfully model emulator



Reduction in PCM writes with KG-N and KG-W versus PCM-Only

Kingsguard collectors limit PCM writes

KG-W much better than KG-N

	Simulation	Emulation
KG-N	4%	8%
KG-W	62%	64%

Increase in execution time with KG-W versus KG-N

KG-W is slower than KG-N because it monitors writes to objects

	Simulation	Emulation
KG-W	+7%	+10%

Graph workload evaluation

GraphChi: Analyze BIG graphs on a single machine
Both Java and **C++** implementations

Page Rank *and* Connected Components

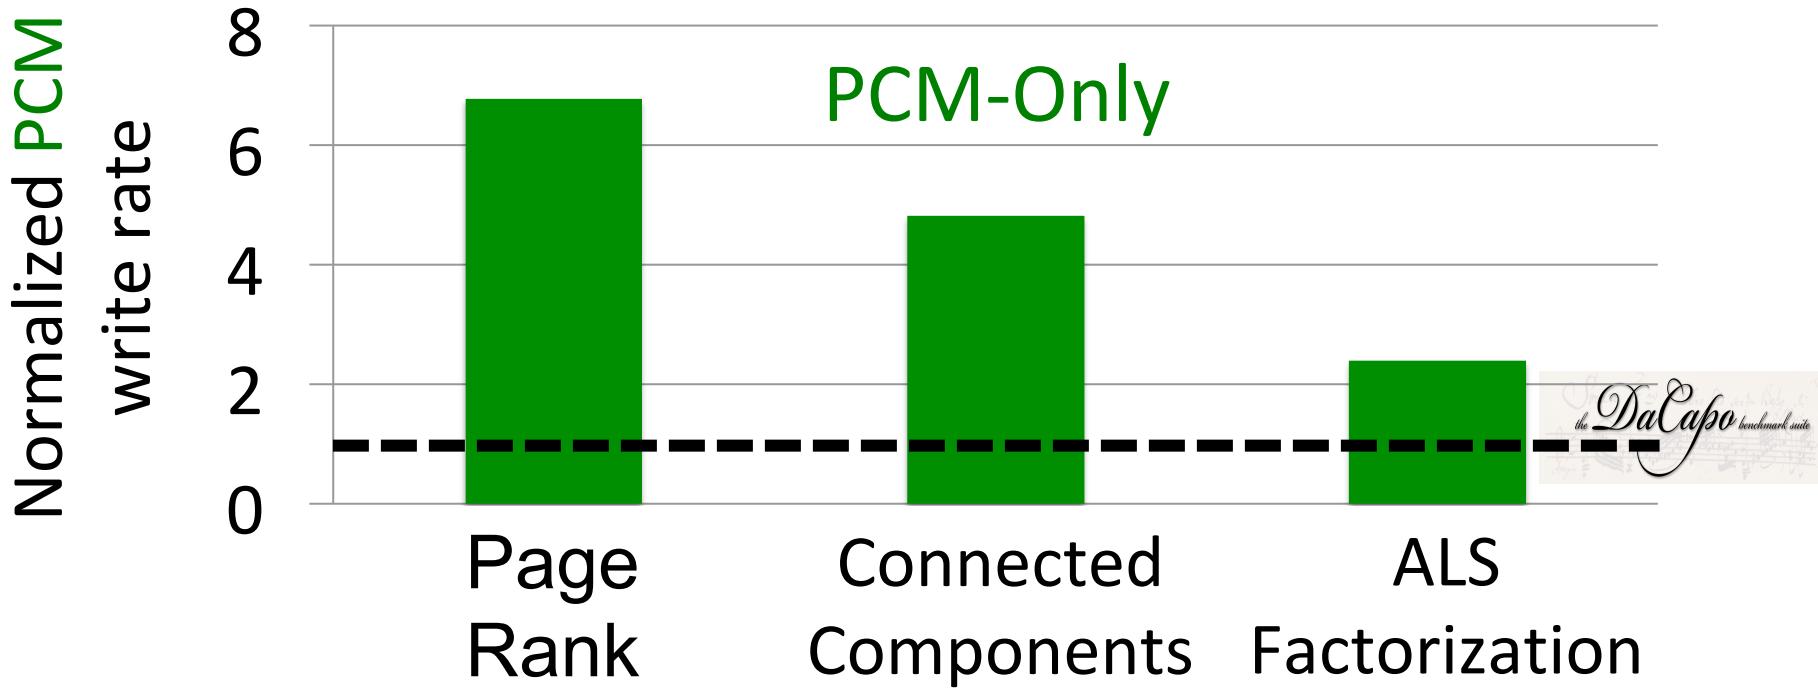
LiveJournal social network

ALS Factorization

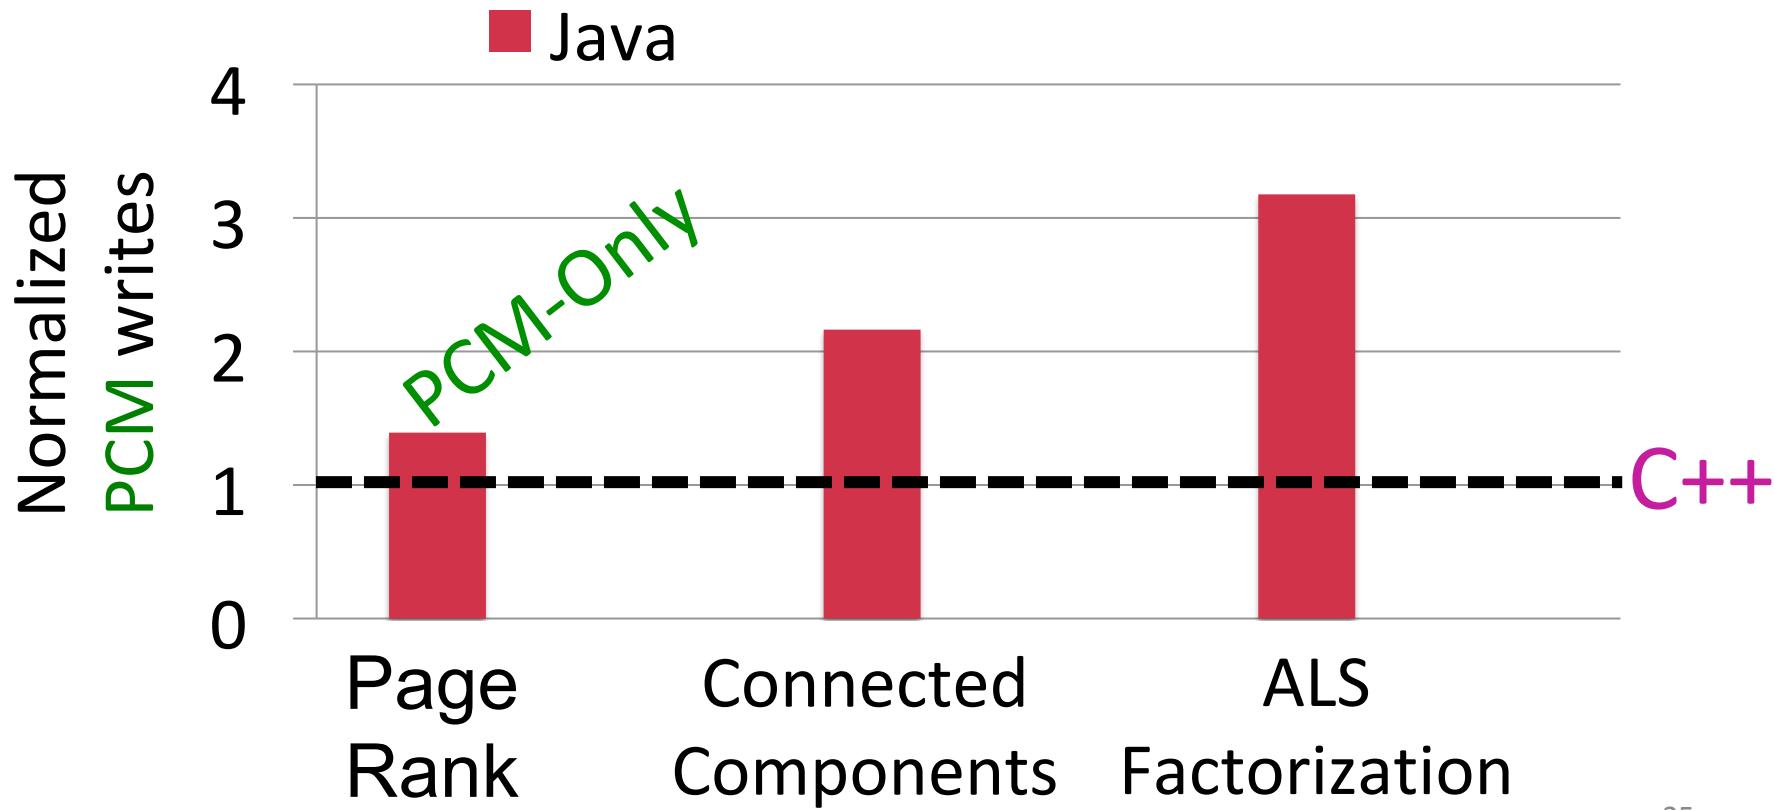
Netflix challenge

Graph apps write more than DaCapo

Billions of vertices → Billions of objects



Java writes more to PCM than C++



Java writes more to PCM than C++

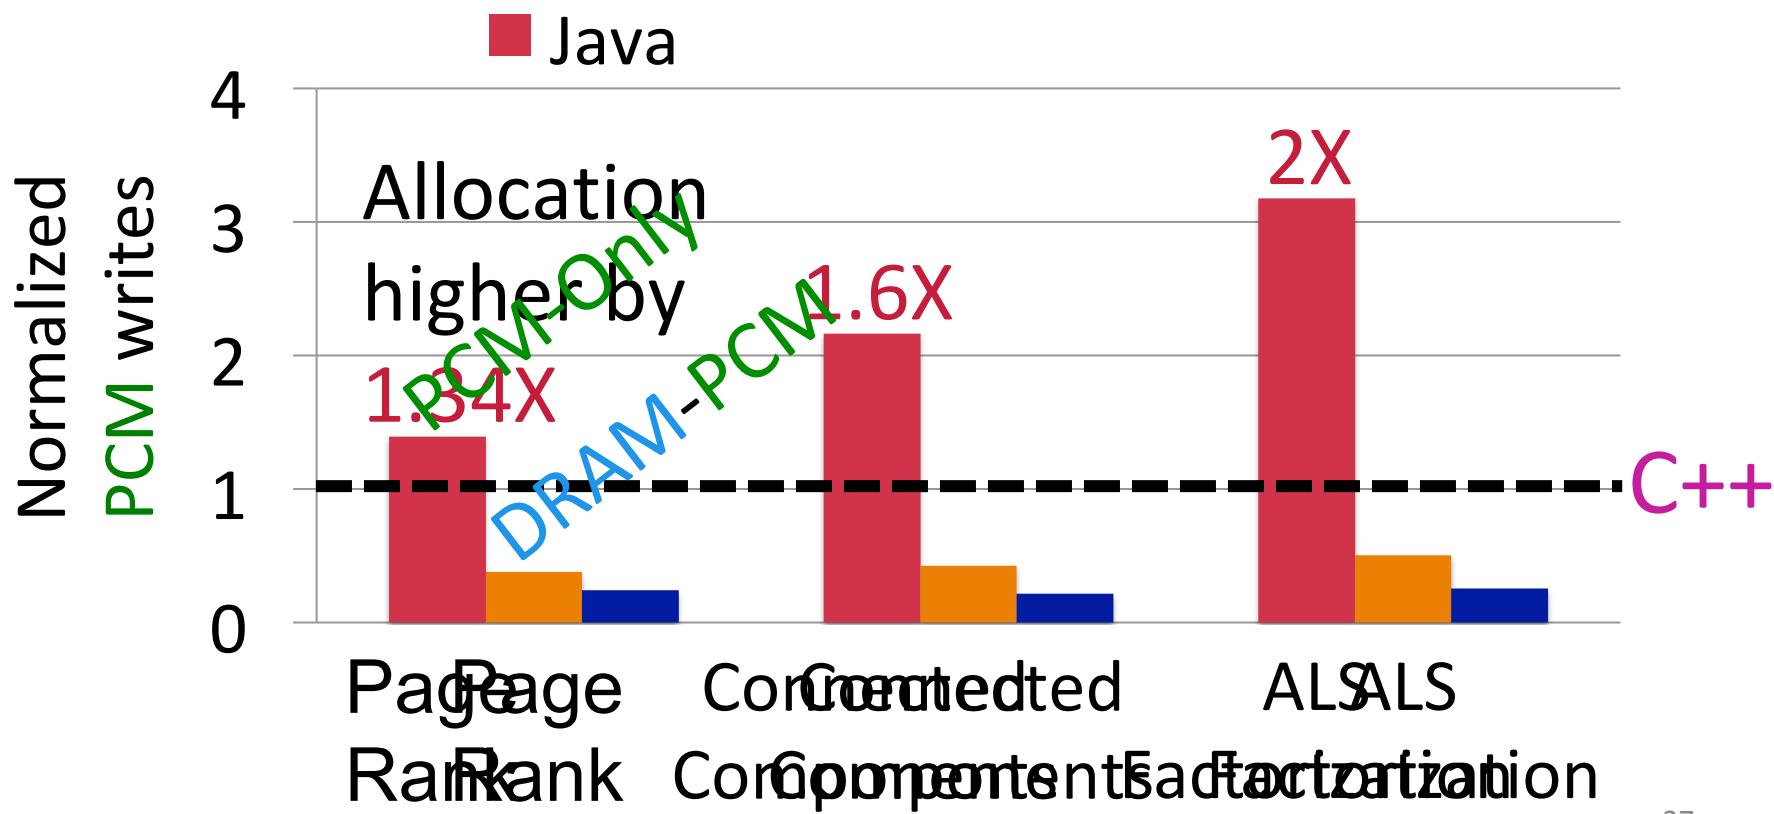
Reasons

Higher allocation volume →

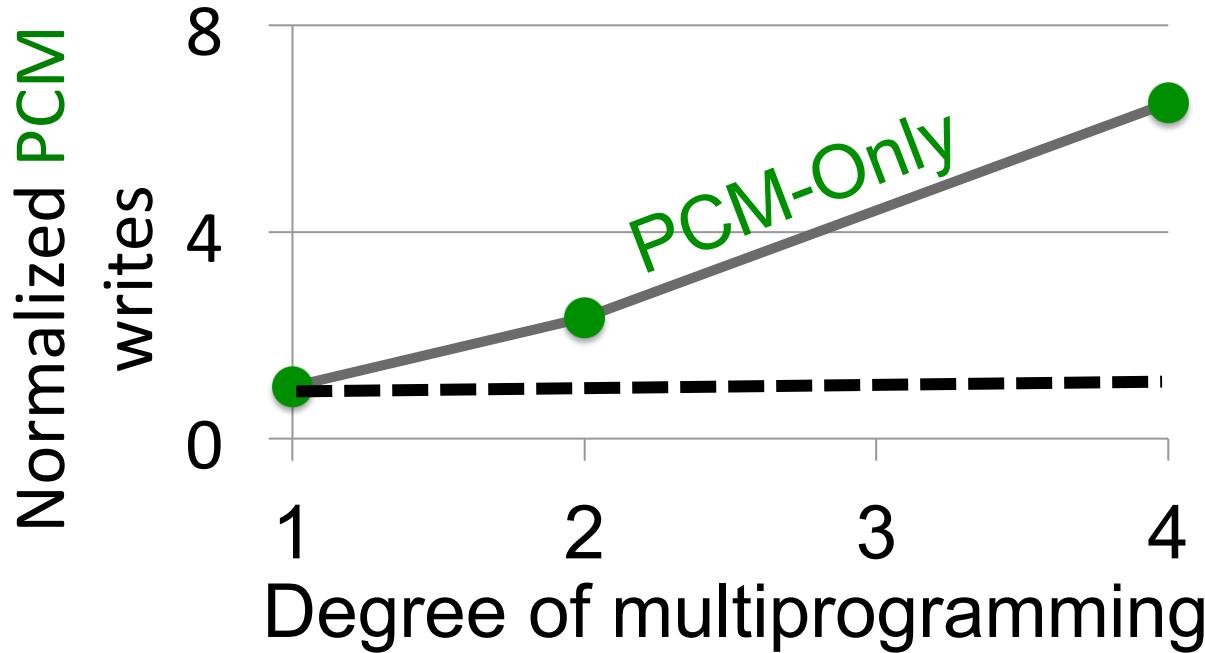
Copying between heap spaces

Zeroing to provide memory safety

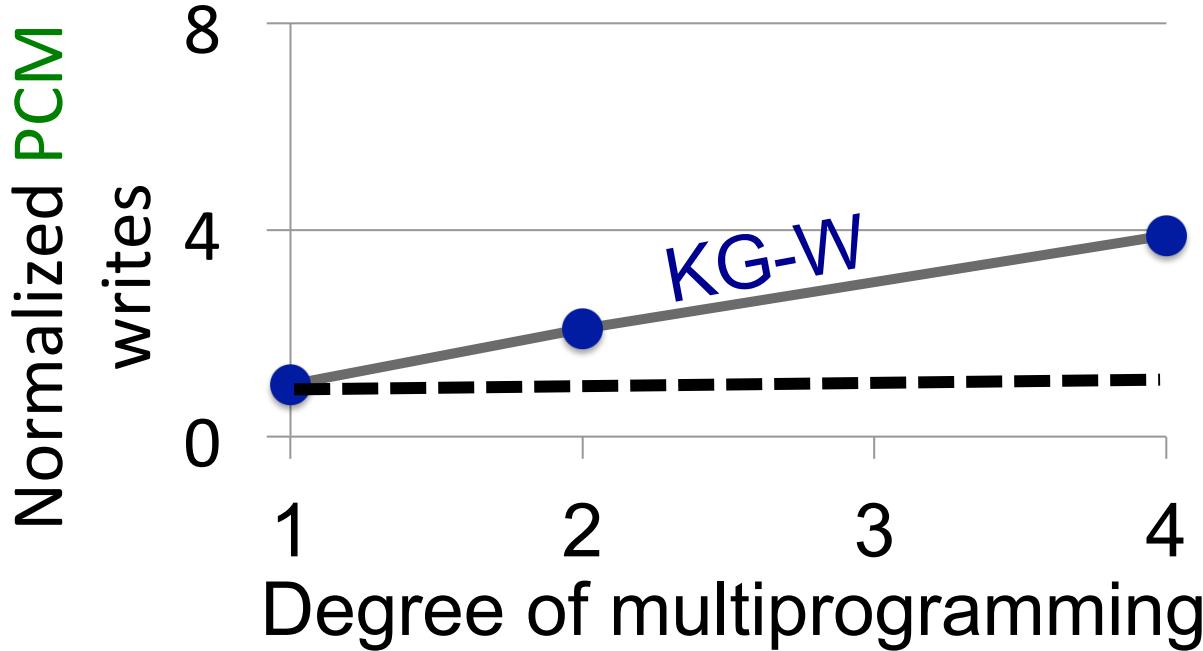
Java writes more to PCM than C++



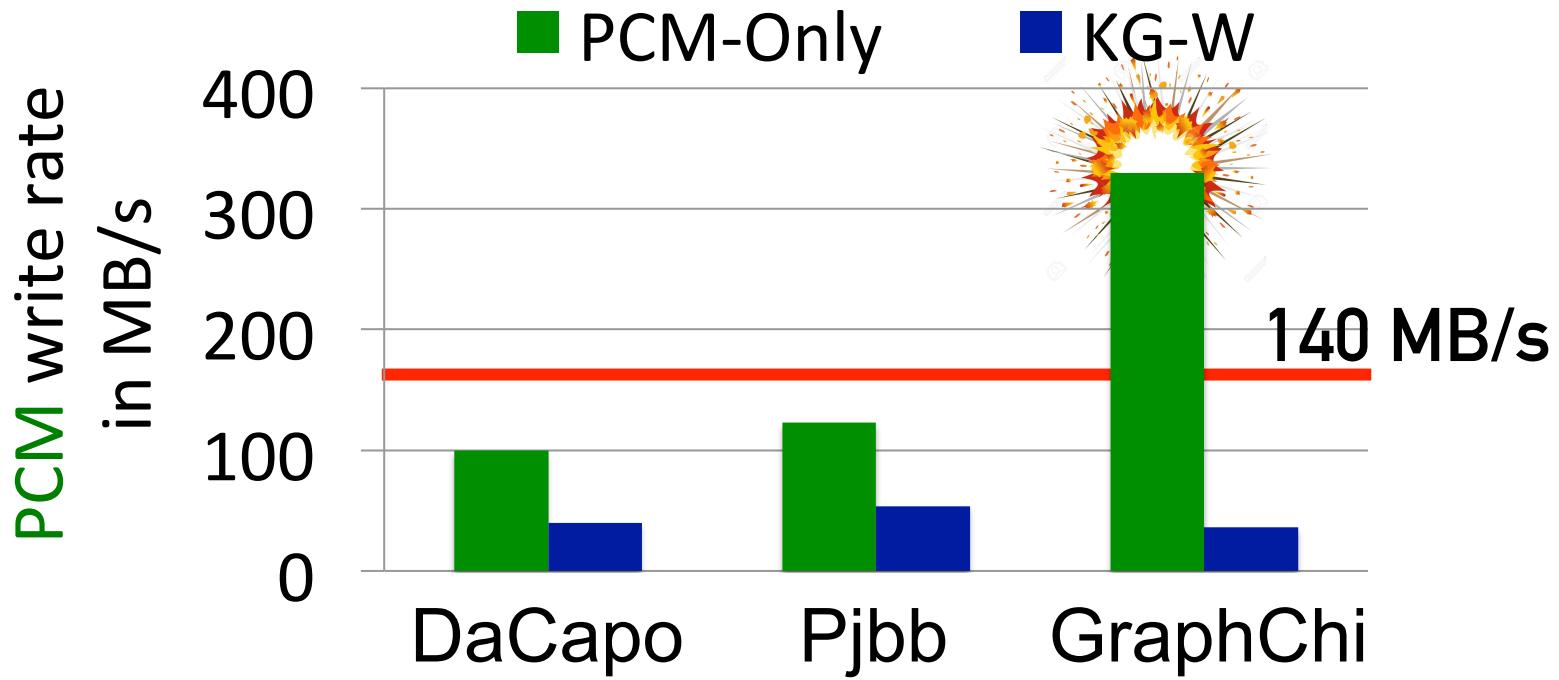
Writes increase **super-linearly** due to multiprogramming with PCM-Only



Writes increase **linearly** due to multiprogramming with KG-W



PCM-Only is not practical as main memory



Conclusions

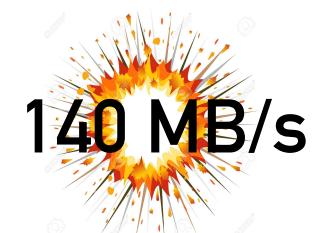
Across the stack emulation of hybrid memory



Similar outcomes with different evaluation methods



More research to make **PCM** practical as main memory



Model	Capacity	DWPD	TBW	GB/Day	Warranty
SSD 600p	1,000GB	0.32	576.0	316	5 Years
SSD 750	1,200GB	0.06	127.0	70	5 Years
DC P3320	1,200GB	0.68	1,480.0	811	5 Years
DC P3520	1,200GB	0.68	1,480.0	811	5 Years
DC D3600	1,000GB	3	5,475.0	3,000	5 Years
DC D3700	1,600GB	10	29,200.0	16,000	5 Years
DC P3500	1,200GB	0.3	657.0	360	5 Years
DC P3600	1,200GB	3	6,570.0	3,600	5 Years
DC P3700	1,600GB	17	49,640.0	27,200	5 Years
DC S3100	1,000GB	0.1	109.5	100	3 Years
Optane™ SSD DC P4800X	375GB	30	12,318.8	11,250	3 Years
Optane™ Memory	16GB	6.25	182.5	100	5 Years
Optane™ Memory	32GB	3.125	182.5	100	5 Years
Source: Objective Analysis, April 2017					