


To expose, or not to expose, hardware heterogeneity to runtimes

Shoaib Akram

Ghent University, Belgium

Shoaib.Akram@UGent.be

Circa 2000: Hardware fixed at design




Out-of-order 1 GHz processor
Volatile DRAM main memory
Persistent disk storage

The shift to power-aware computing

Dennard scaling has stopped


Constant power density as transistor size shrinks

Reliance on battery-operated devices




Heterogeneous Multicores

Each core has its own frequency domain (DVFS)



DRAM price and supply trends



Hybrid DRAM-PCM memory

- 😊 More GB/\$ with Phase Change Memory
- 😢 Higher latency *and low* endurance

Speed
Endurance

Capacity
Persistence

DRAM

PCM

Some challenges of heterogeneity

Schedule threads on **big-small** cores

Regulate **DVFS**

Mitigate **PCM** wear-out

Bridge **DRAM-PCM** latency-gap


...

Hardware/OS only approaches

Hardware exposes counters

OS predicts how software behaves


OS configures knobs and manages heterogeneity



Pros/cons of moving up the stack

Gain in semantic
knowledge 😊

Loss in abstraction 😥



Exposing heterogeneity to runtimes

😊 Proactive

Thread X chases pointers (memory-bound)

$X \rightarrow Y$ is producer \rightarrow consumer

Memory region X is highly written

😊 Flexible granularity

Memory mgmt in OS fixed at page granularity

Exposing heterogeneity to runtimes

⌚ Dependency

Hardware vendor relies on Microsoft/Oracle etc
OS is ubiquitous

⌚ Software complexity

Gain insight into software behavior
Design, code, verify

⌚ Native applications

C has a non-negligible fan base

Beyond heterogeneity

Hardware multithreading

Turbo boost

Prefetching

Variable page sizes

Cache and memory-bandwidth partitioning

Accelerators and FPGAs

3D Stacked memory

Outline

Garbage collection for hybrid memories

Concurrent collection on heterogeneous multicores

Managing DRAM-PCM memory

Mitigate PCM wear-out ✓

Bridge the DRAM-PCM latency gap

Speed
Endurance

Capacity

DRAM


PCM

Managing DRAM-PCM memory



*Write-Rationing
Garbage Collection
for Hybrid Memory*

Operating System
Coarse-grained
pages  KB



Garbage collection
Proactive ☺
Fine-grained
objects ● ● ● ●

GC manages DRAM-PCM hybrid better than OS

DRAM heap management

Heap Tracker



available



occupied



HEAP_BEGIN

HEAP_END

Heap Organization



DRAM heap management

Heap Tracker

✓ available


★ occupied




HEAP_BEGIN

HEAP_END

Heap Organization



Physical Memory



DRAM-PCM heap management

Heap Tracker



available



occupied



DRAM_BEGIN



PCM_BEGIN

PCM_END

Heap Organization



'mbind("pcm")'



Physical Memory

Kingsguard-Nursery (KG-N)



Write-rationing GC: concentrate writes in DRAM







70%

of writes



22%

to 2% of objects










Kingsguard-Writers (KG-W)



KG-W monitors writes in a DRAM observer space

Trades off performance for better endurance



More optimizations in KG-W



Short-lived large objects in DRAM

Large data-structures cause writes to PCM

Keep them in DRAM

Object meta-data in DRAM


GC updates to mark bits lead to writes to PCM

Keep them in DRAM

OS versus Kingsguard



DRAM




PCM

Rank pages according to writes

A page with \mathcal{T} writes is a DRAM candidate

Adjust for phase behavior


OS versus Kingsguard



Average for 7 DaCapo benchmarks in simulation

KG-W reduces 3X more writes than OS-Only

OS versus Kingsguard




Average for 7 DaCapo benchmarks in simulation

OS predicts nursery pages but migrations harmful

Emulation on NUMA hardware



PCM-Only is not practical as main memory




Crystal Gazer: Profile-Driven Write-Rationing Garbage Collection for Hybrid Memories



Heterogeneous multicore scheduling

Mutator → big, Concurrent GC → **big** or **small**?




GC on big or small




Slow GC pauses application due to no memory


GC on big or small



GC on big or small



GC on big or small




GC-criticality-aware scheduling


Runtime detects GC-**criticality**

Communicates **criticality** to the OS

OS adjusts GC **priority**



Better energy efficiency with GC-criticality-aware scheduling



To expose, or not to expose, hardware heterogeneity to runtimes

Always need OS (supervisory role)

Virtual memory, thread migration, and so on

Language runtimes can guide OS in forming
the best policies to manage emerging hardware