

# Fairness-Aware Scheduling on Single-ISA Heterogeneous Multi-Cores

**Kenzo Van Craeynest<sup>+</sup>**

Shoaib Akram<sup>+</sup>

Wim Heirman<sup>+</sup>

Aamer Jaleel\*

Lieven Eeckhout<sup>+</sup>

+ Ghent University

\* VSSAD, Intel Corporation

# Single-ISA heterogeneous multi-cores

## Multiple core types

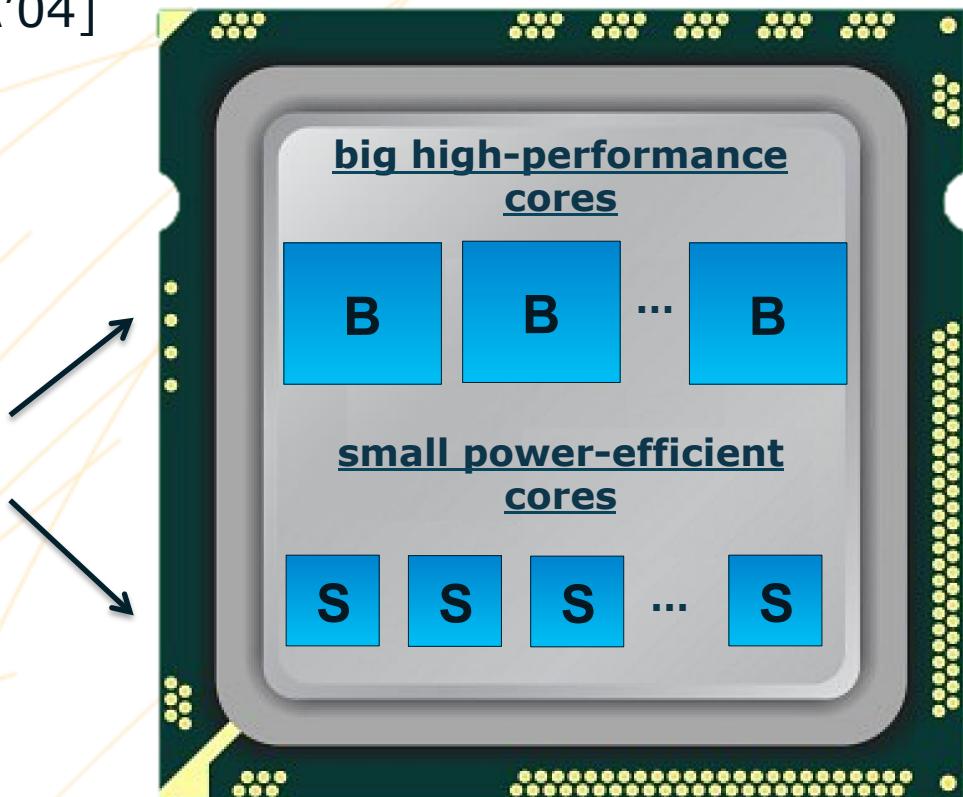
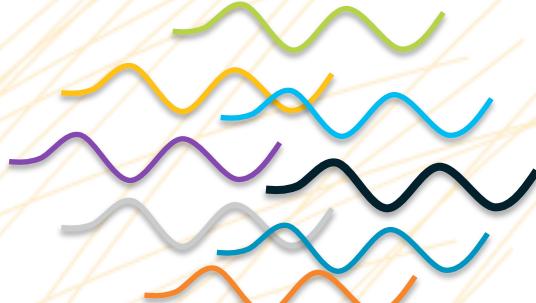
- representing different power/performance trade-offs

## Well-established power benefits

- [Kumar et al. MICRO'03, ISCA'04]

## Commercial examples

- Big.LITTLE, Kal-El



# Prior Work: Put the Thread That Will Benefit the Most on the Big Core

Many different scheduling techniques

- Static scheduling
- Sampling-based scheduling

*Kumar et al., ISCA'04; Patsilaras et al., TACO'12*

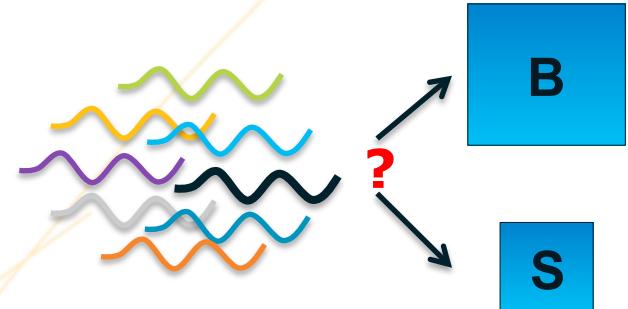
- Proxies for performance

*Memory-domance (Becchi et al., JILP'08; Koufaty et al., EuroSys'10;  
Sheleпов et al., OS Review'09)*

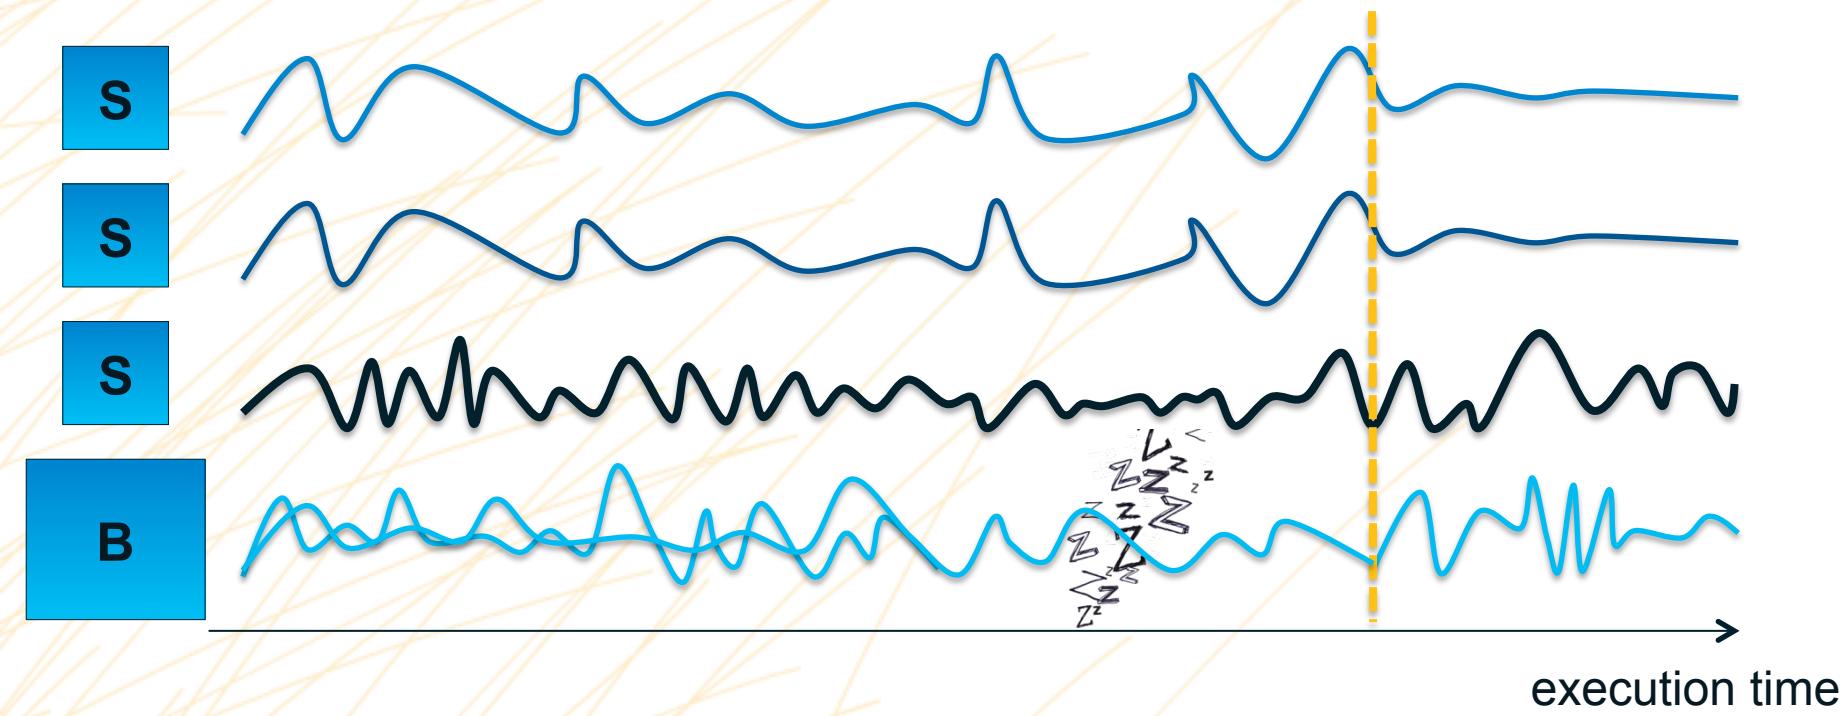
*Age-based Scheduling (Lakshminarayana et al., SC'09)*

- Model-based scheduling

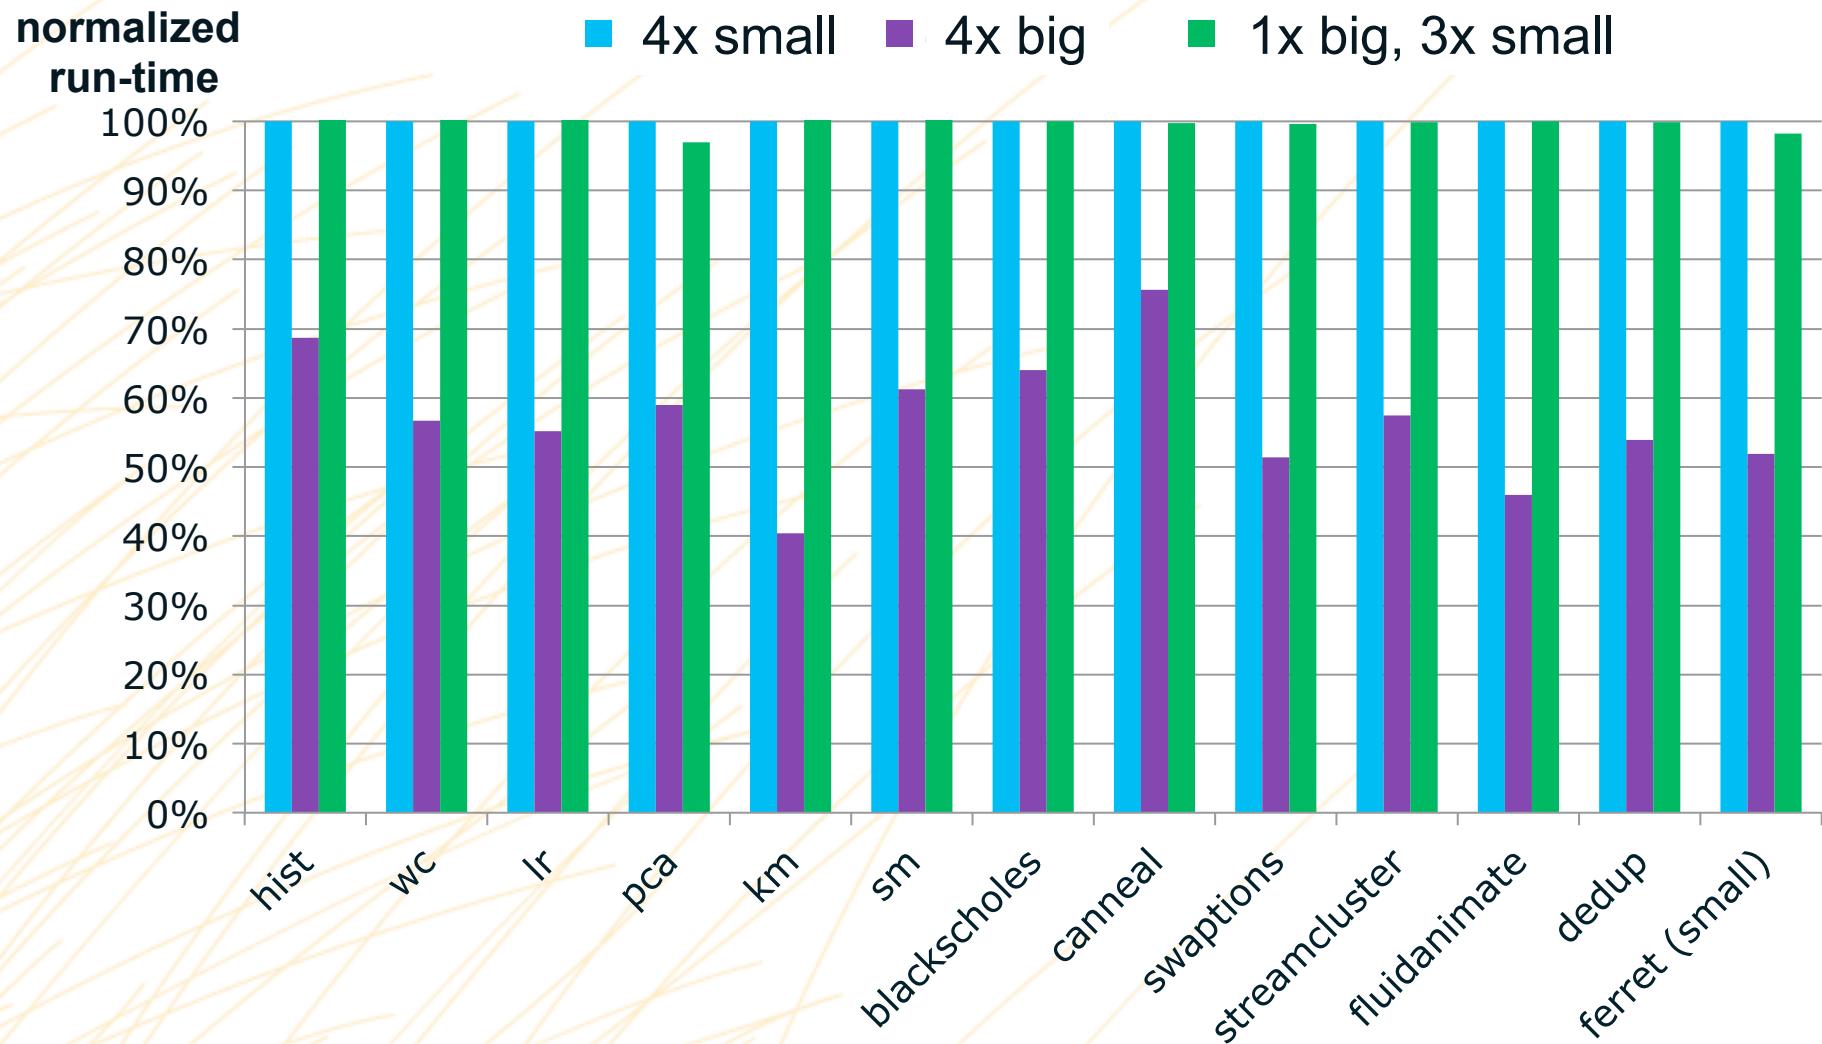
*Van Craeynest et al., ISCA'12; Lukefahr et al., MICRO'12*



# Traditional Scheduling can be Suboptimal



# Threads pinned on Small Cores Determine Performance



# Fairness-Aware Scheduling on Single-ISA Heterogeneous Multi-Cores

Scheduling methodologies that aim to improve fairness

- *Equal-time scheduling*
- *Equal-progress scheduling*

Will show that Fairness-Aware Scheduling

- *Significantly improves fairness*
  - *Allowing QoS, accounting,...*
- *Significantly reduced run-time for many multi-threaded applications over state-of-the-art throughput-optimizing scheduling*

# Fairness for Heterogeneous Multi-Cores

$$\text{slowdown} = S \downarrow i = T \downarrow \text{het}, i / T \downarrow \text{big}, i$$

Number of cycles to execute a thread on a heterogeneous multi-core

Number of cycles to execute a thread in isolation on big core

**Schedule is fair if slowdown of all running threads is the same**

$$\text{fairness} = 1 - c \downarrow S = 1 - \sigma \downarrow S / \mu \downarrow S = 1 - \text{std\_dev}(S) / \text{avg}(S)$$

Coefficient of variation, a measure of unfairness

# Experimental Setup



Simulated hardware

	<b>small</b>	<b>big</b>
<b>issue width</b>		4-wide
<b>clock frequency</b>		2.6 GHz
<b>cache hierarchy</b>	32KB (p) / 256 KB (p)/	16MB (s)
<b>μarch</b>	in-order	out-of-order

Sniper:

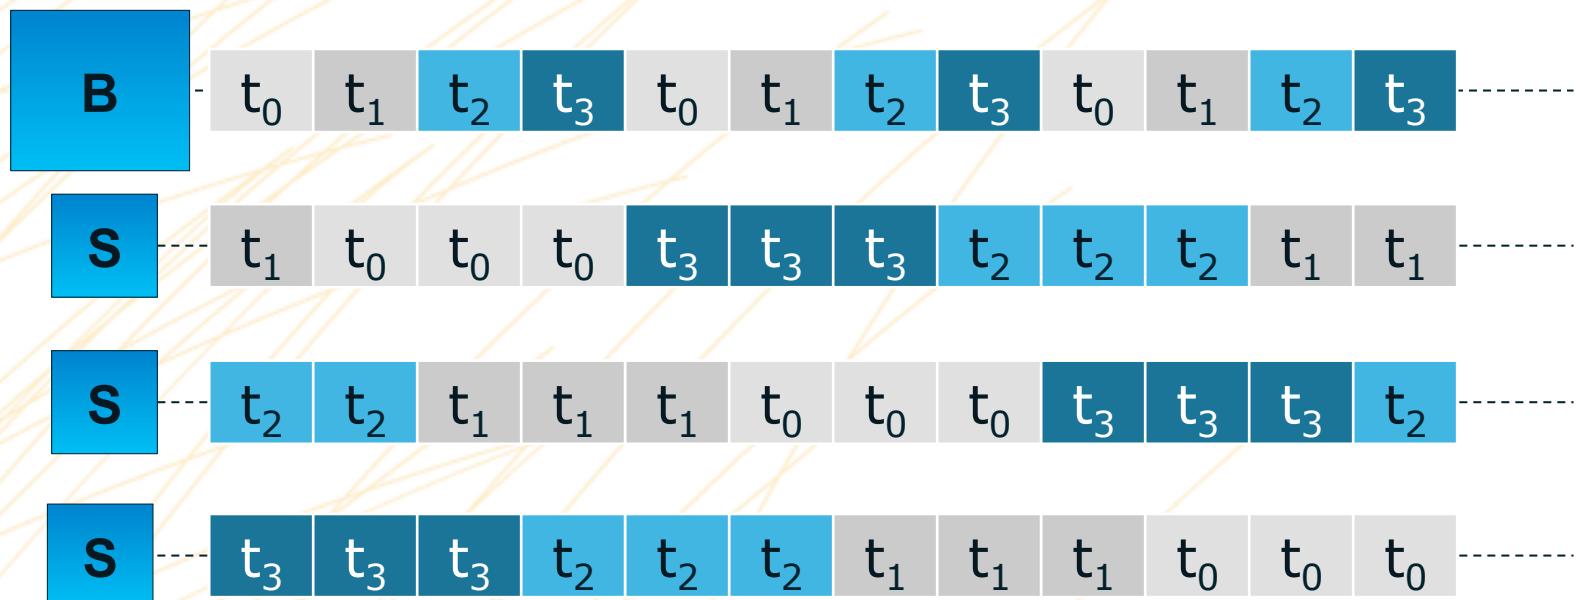
- parallel, hardware-validated x86-64 multi-core simulator

Multi-threaded and multi-programmed workloads

- spec2006, PARSEC and MapReduce

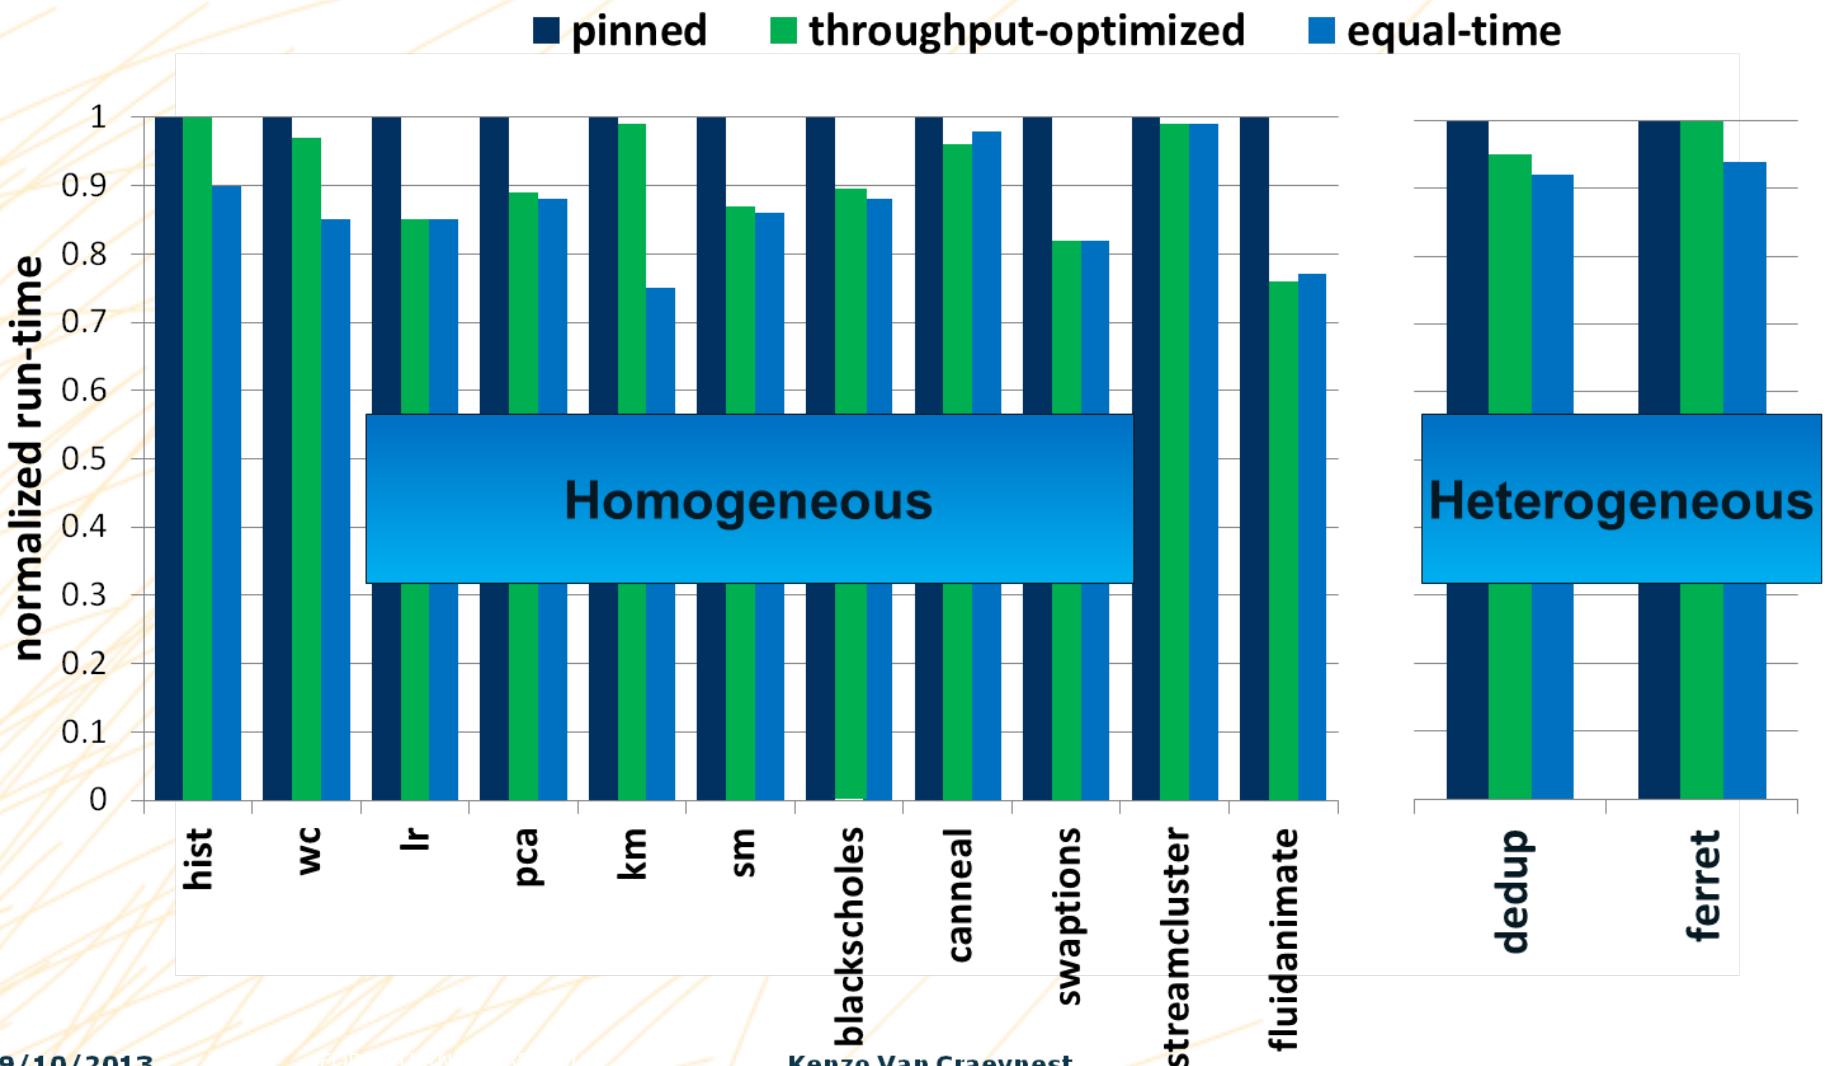
# Achieving Fairness: Equal-time Scheduling

- Each thread runs for same amount of time on each core type
- Can be implemented with minor changes to a Round-robin scheduler



# Optimizing for Fairness Reduces Run-time for Homogeneous Multi-Threaded Workloads

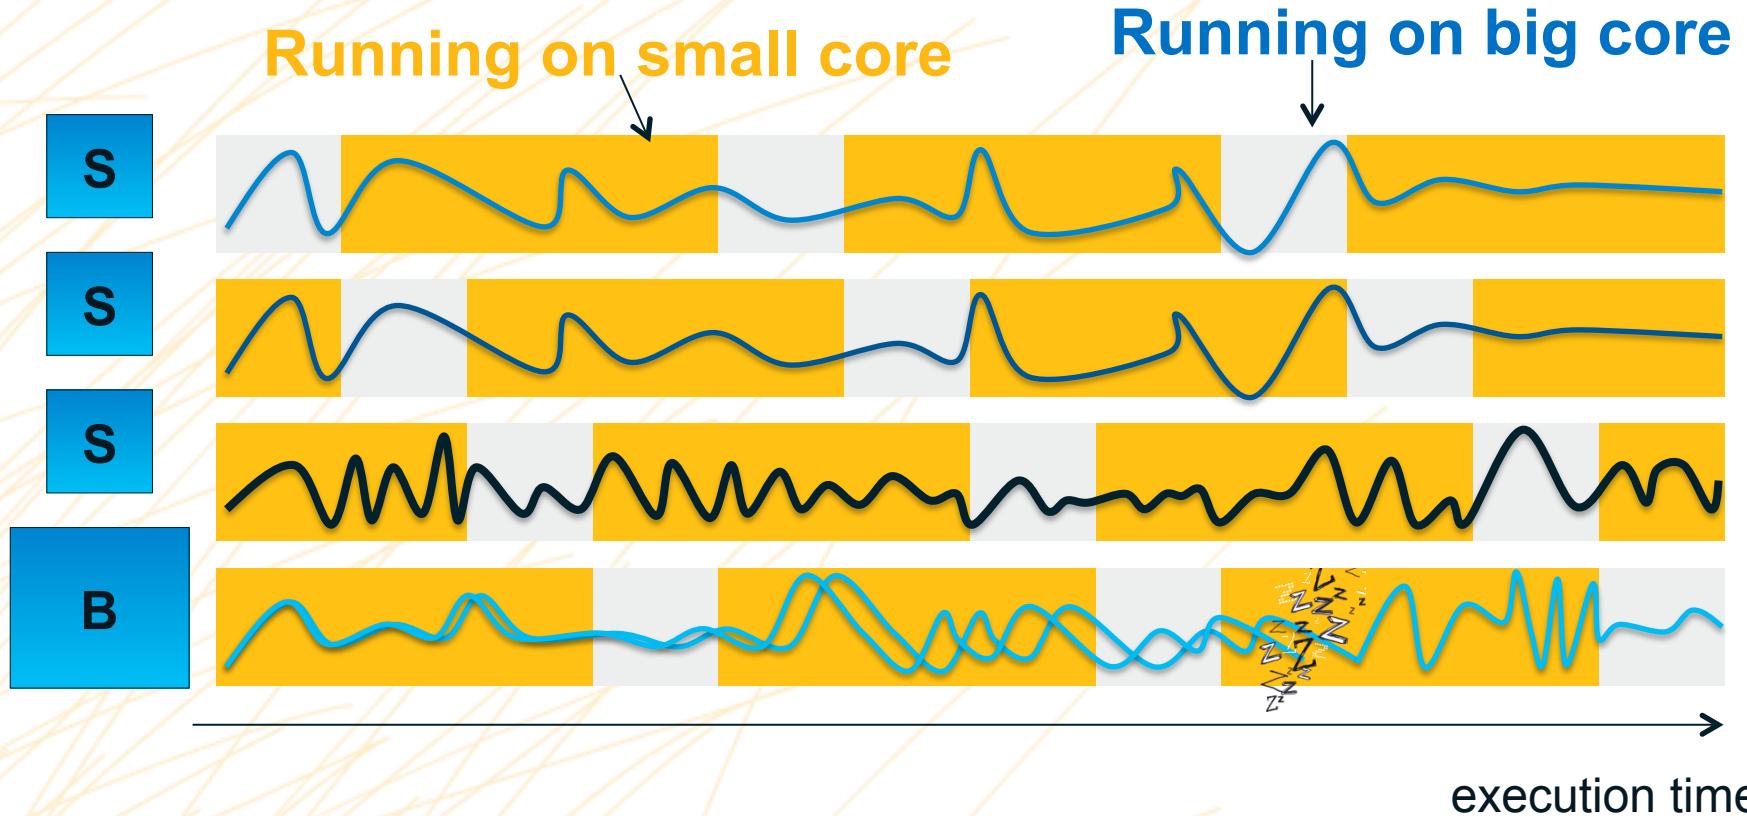
1B3S system



# Equal-Time Doesn't Guarantee Equal-Progress

Some threads experience a larger slowdown than others

- Equal time on different core types ≠ equal progress
- Therefore fairness is not guaranteed



# Achieving Fairness: Equal-progress Fairness-Aware Scheduling

- Guarantee that all threads make the same progress compared to their big-core performance
- Continuously monitor fairness and adjust schedule to achieve fairness

$$S_{\downarrow i} = T_{\downarrow het,i} / T_{\downarrow big,i} = T_{\downarrow big,i} + T_{\downarrow small,i} / T_{\downarrow big,i} + T_{\downarrow small,i} / R_{\downarrow i}$$

Overall slowdown of  
the thread

Scale execution time  
on small core

Performance ratio between  
big and small core

# Estimating the Performance Ratio

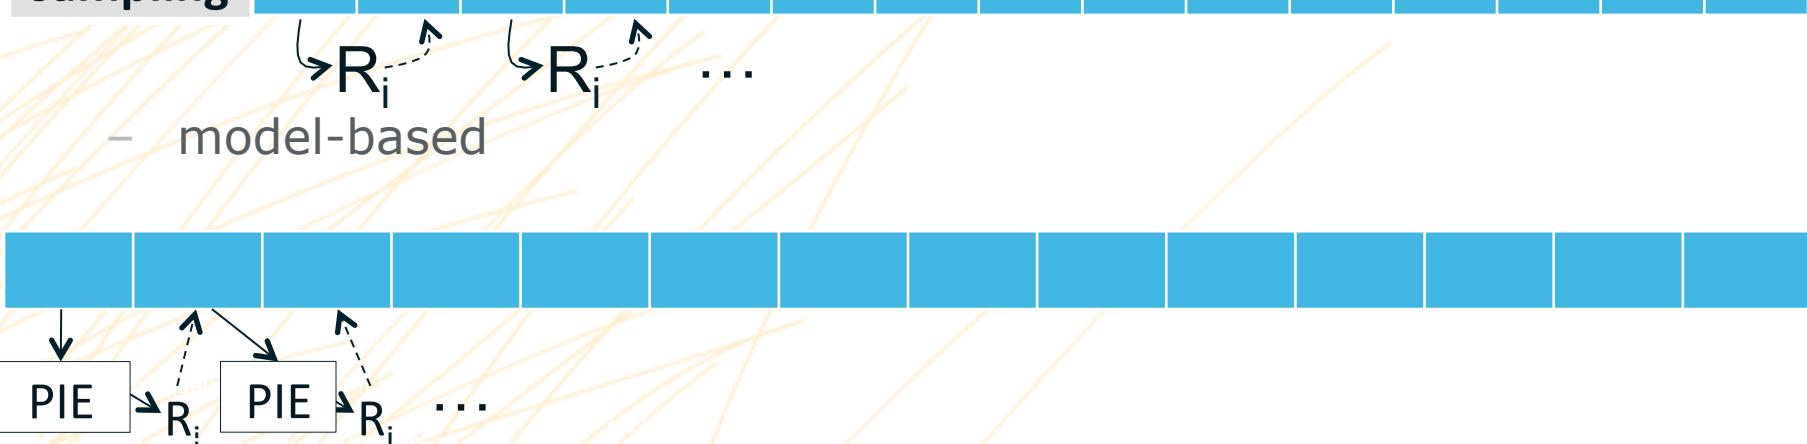
- Proposed 3 methods
  - sampling-based



- history-based



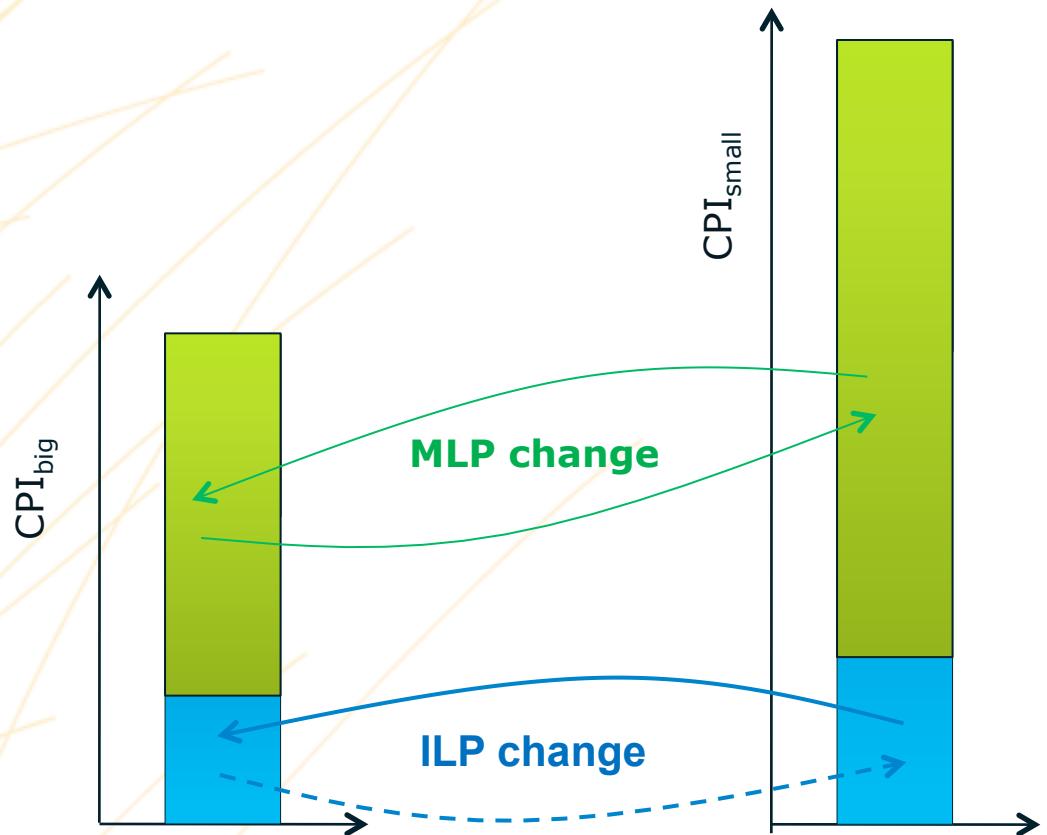
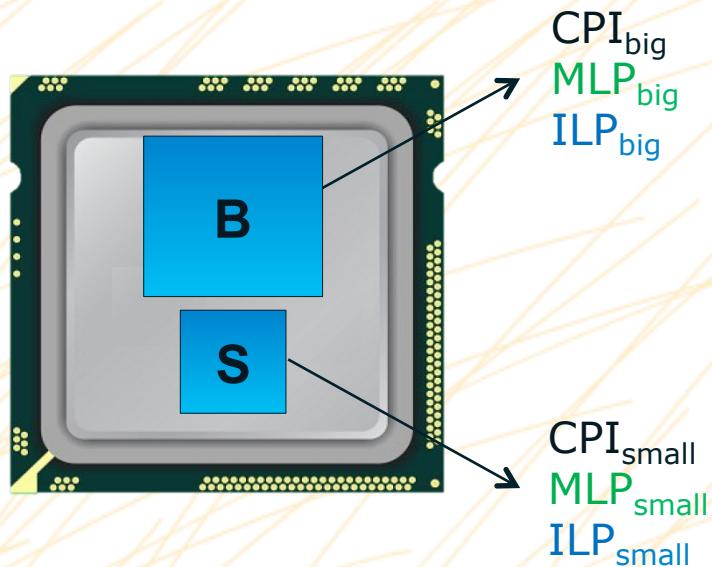
- model-based



# Performance Impact Estimation (PIE)

[Van Craeynest et al., ISCA'12]

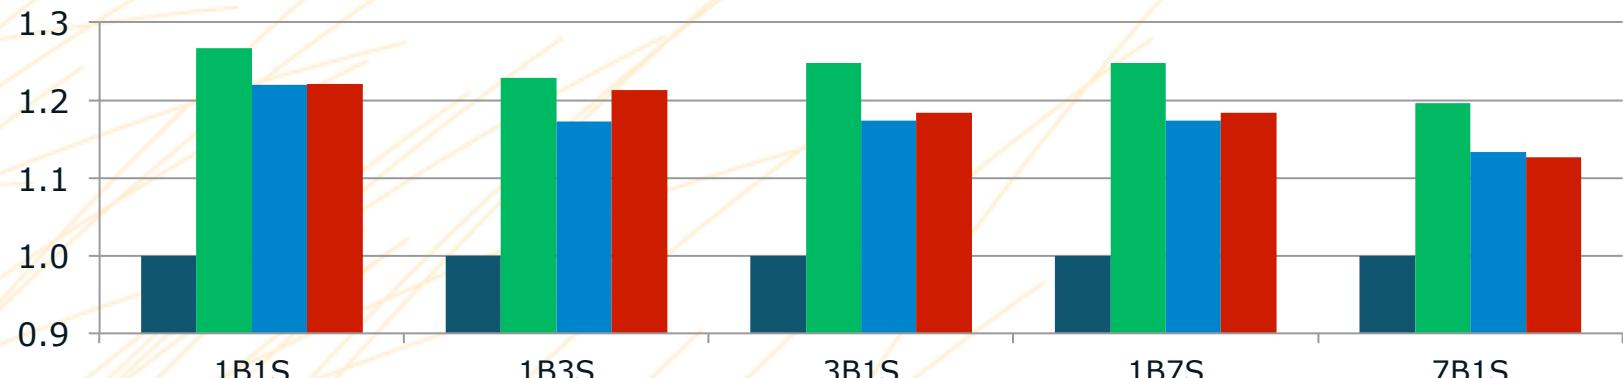
1. Determine where application spends its execution time
2. Use change in MLP exposed to predict change in  $\text{CPI}_{\text{mem}}$
3. Use change in ILP exposed to predict change in  $\text{CPI}_{\text{base}}$



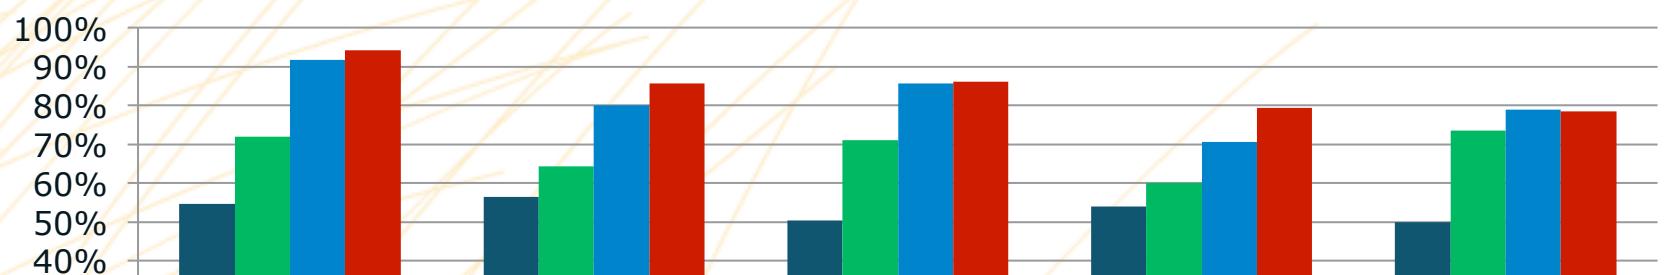
# Fairness-aware Scheduling Across Configurations for Multi-Programmed Workloads

■ pinned ■ throughput-optimized ■ equal-time ■ equal-progress

normalized throughput



fairness

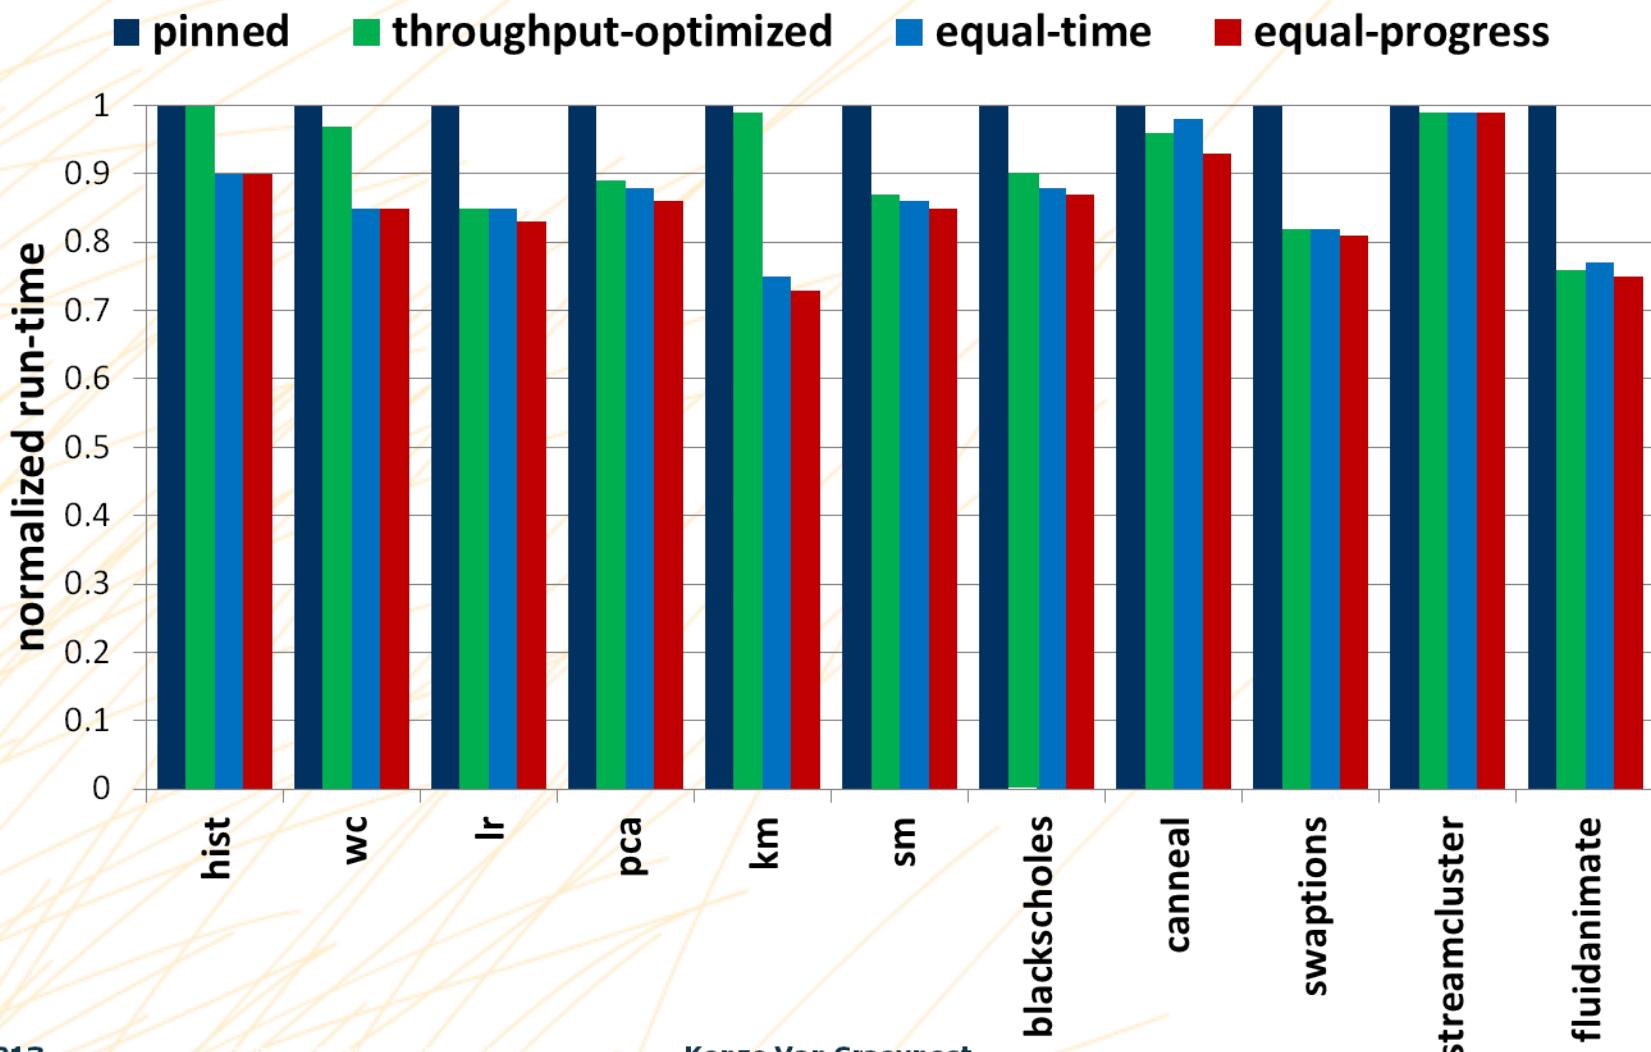


QoS, cycle-accounting , abstraction of heterogeneity,...

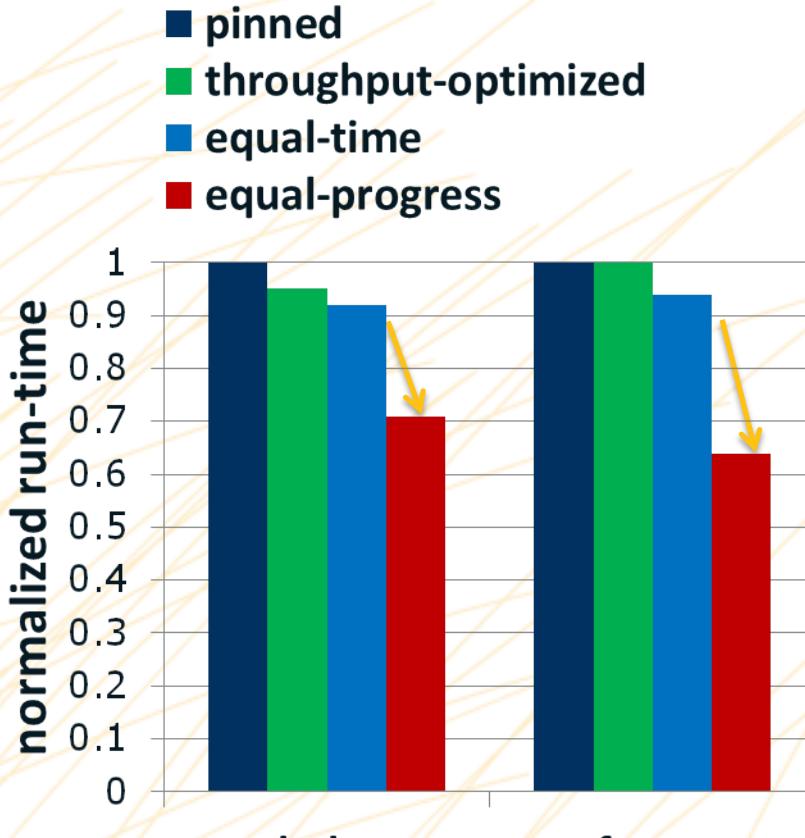


# Optimizing Fairness Reduces Run-time for Homogeneous Multi-Threaded Workloads

1B3S system



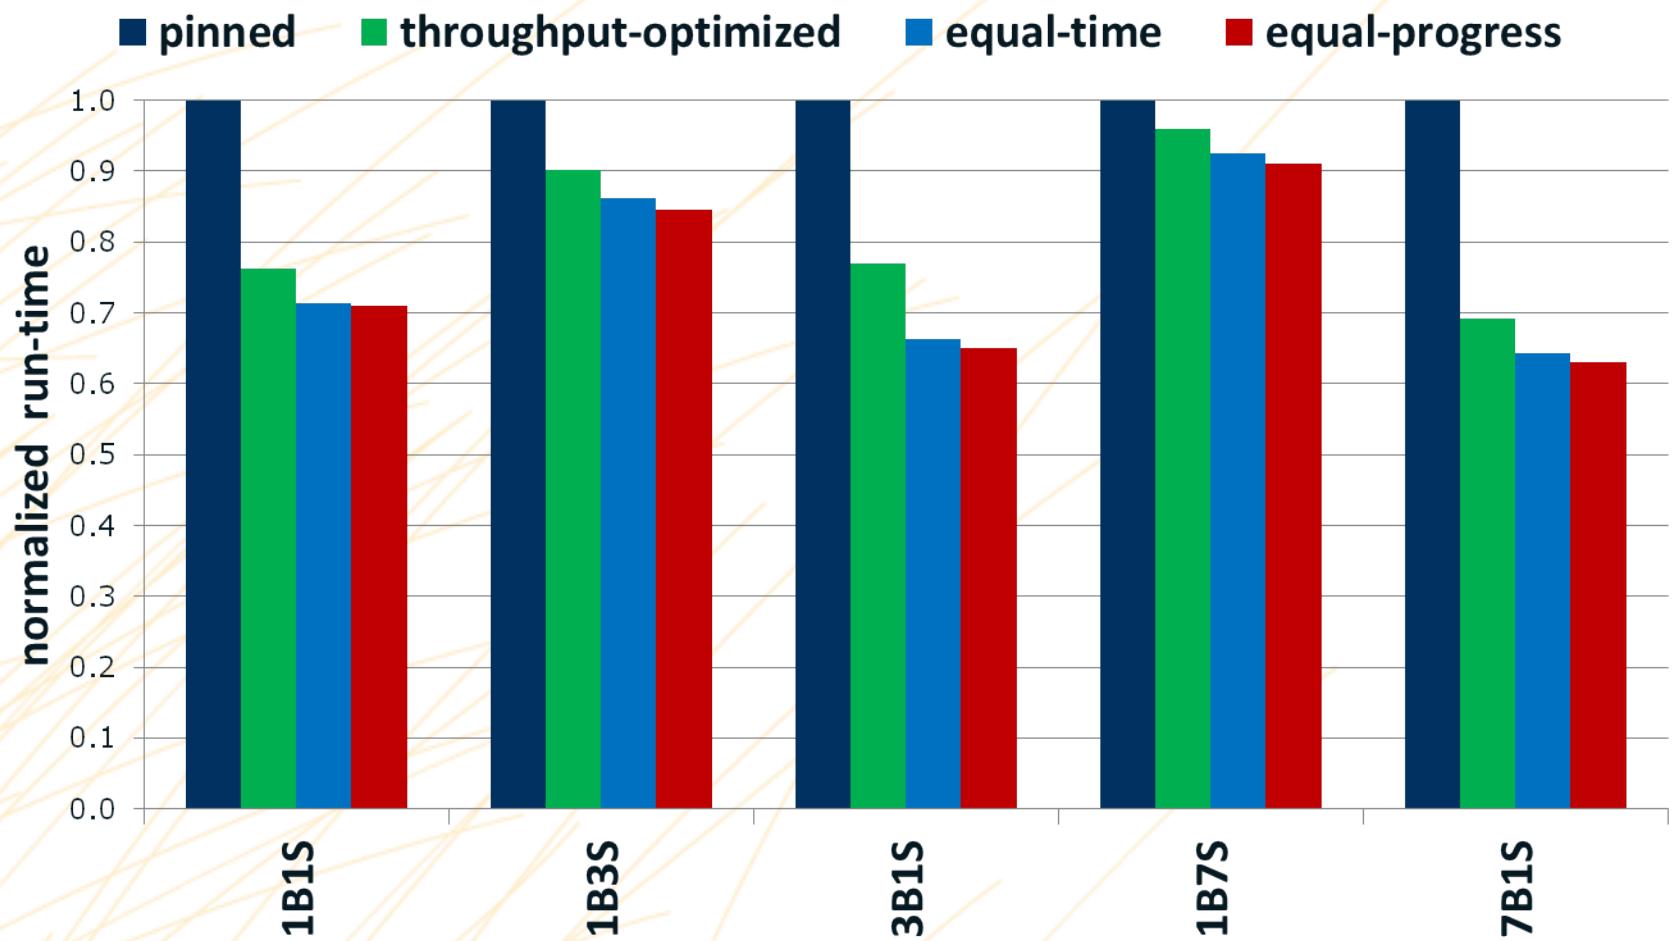
# Optimizing for Fairness Reduces Run-time for Heterogeneous Multi-Threaded Workloads



1B3S

- Heterogeneous applications
  - Threads can have different performance ratio
- Equal-time scheduling does not result in a fair schedule
- Equal progress greatly reduces run-time over throughput-optimized AND equal-time scheduling for heterogeneous multi-threaded applications

# Fairness-aware Scheduling Across Configurations for Homogeneous Multi-Threaded Workloads



# Conclusions and Contributions

## Proposed Fairness-optimizing scheduling

- Two methods: equal-time and equal-progress

## Multi-program workloads

- Achieves average fairness of 86% for a 1B3S system while within 3.6% performance of throughput-optimizing scheduling
- Allows for QoS, cycle-accounting, etc. in heterogeneous systems

## Multi-threaded workloads

- Unfair performance results in no performance benefits from heterogeneity
  - Threads running on a big core wait at barriers for threads running on small core
- Average 14% (and up to 25%) performance improvement over pinned scheduling



# Questions?