

# Crystal Gazer: Profile-Driven Write-Rationing Garbage Collection for Hybrid Memories

**Shoaib Akram (Ghent), Jennifer B. Sartor (Ghent and VUB),  
Kathryn S. McKinley (Google), and Lieven Eeckhout (Ghent)**

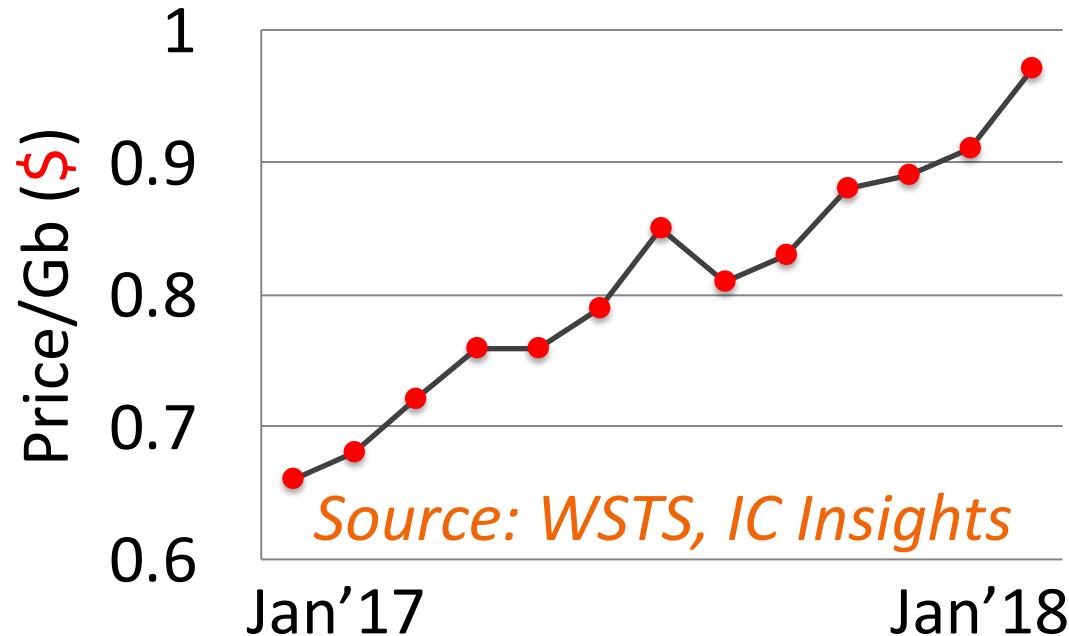
**[Shoaib.Akram@UGent.be](mailto:Shoaib.Akram@UGent.be)**



# Main memory capacity expansion

DRAM → Charge storage a scaling limitation

*Manufacturing  
complexity makes  
DRAM pricing  
volatile*



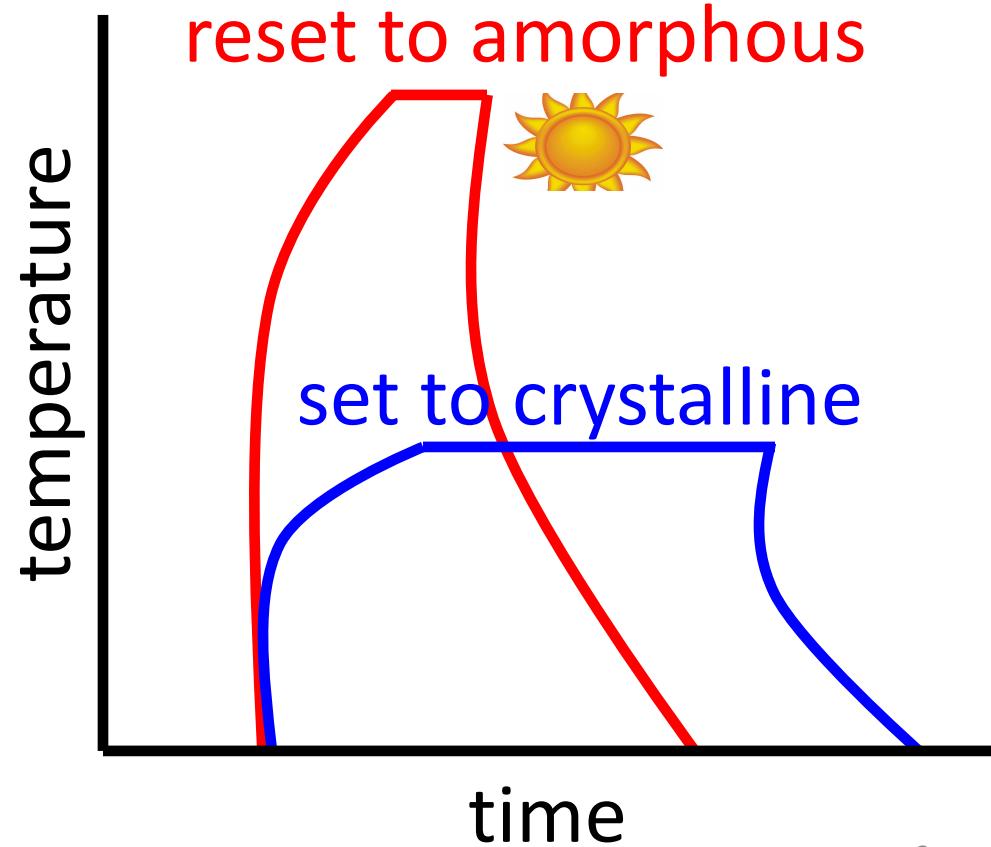
# Phase change memory (PCM)

More Gb/\$

Byte addressable

Latency → DRAM

 Write endurance



# Hybrid DRAM-PCM memory

Speed  
Endurance

Capacity

---

DRAM

---

PCM

PCM alone can wear out in a few months time

*This work → Use DRAM to limit PCM writes*

# Garbage Collection to limit PCM writes

GC understands memory semantics

A GC approach is *finer grained* than OS approaches

Application

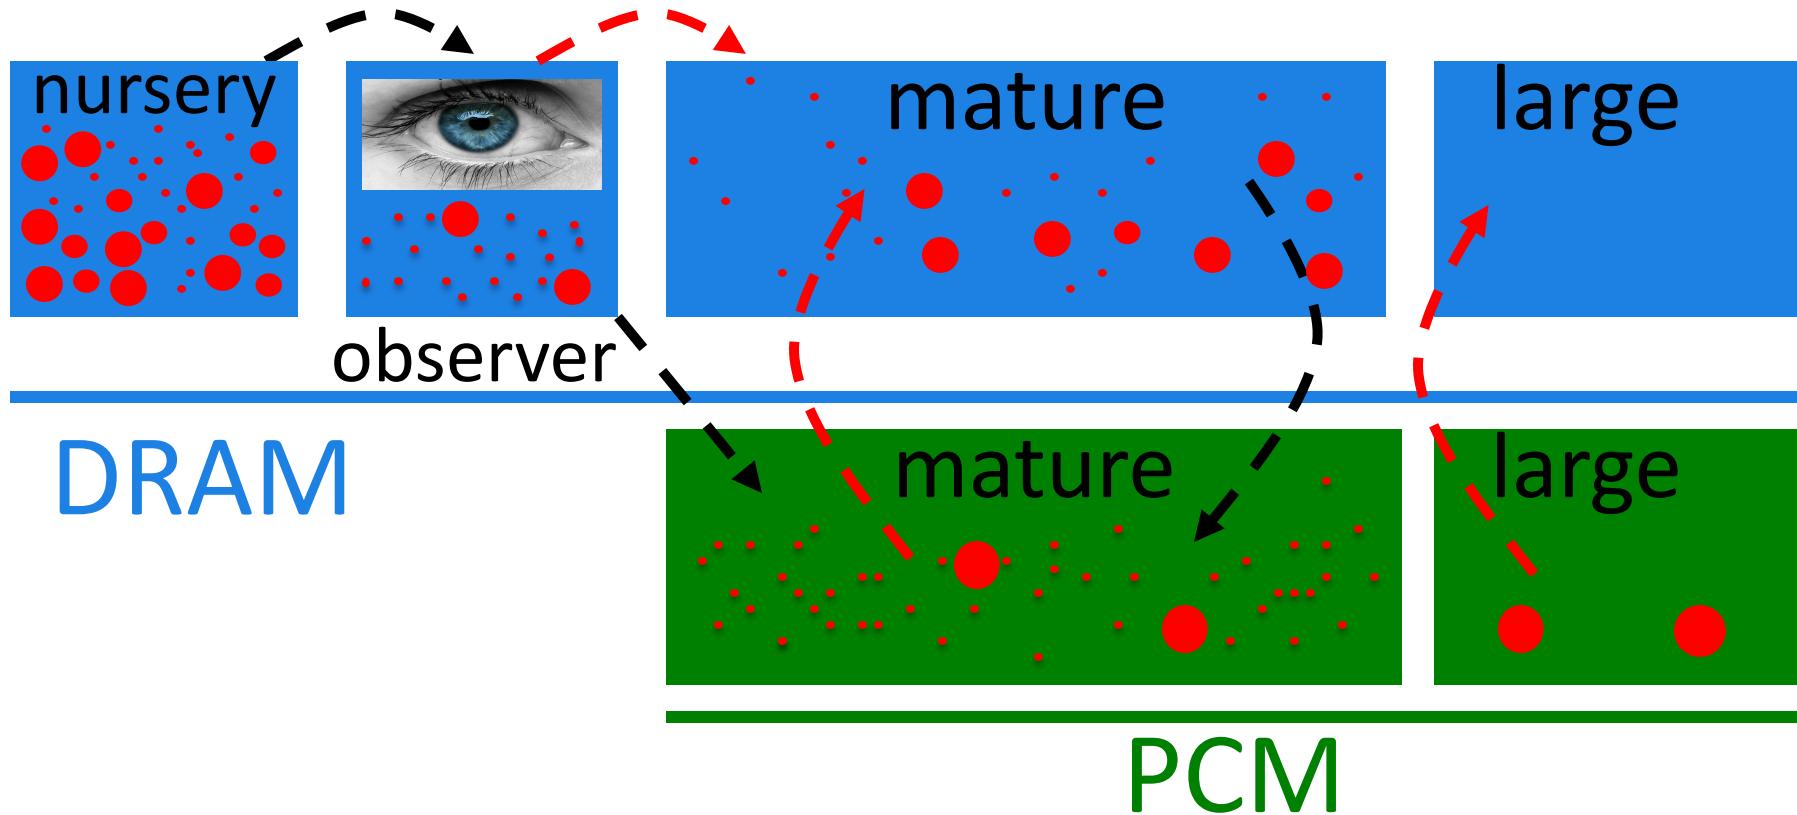


Operating System

Hardware

*Write-Rationing Garbage Collection for Hybrid Memories, PLDI, 2018*

# KG-W Kingsguard-Writers



# KG-W drawbacks

Overhead of dynamic monitoring

Limited time window to predict write intensity  
→ mispredictions

Excessive & fixed DRAM consumption

# Predicting highly written objects without a **DRAM** observer

## Crystal Gazer



# Allocation site as a write predictor

```
a = new Object()  
b = new Object()  
c = new Object()  
d = new Object()
```

# Allocation site as a write predictor

```
a = new Object()
b = new Object()
c = new Object()
d = new Object()
```

Uniform distribution 😞

# Allocation site as a write predictor

```
a = new Object()  
b = new Object()  
c = new Object()  
d = new Object()
```

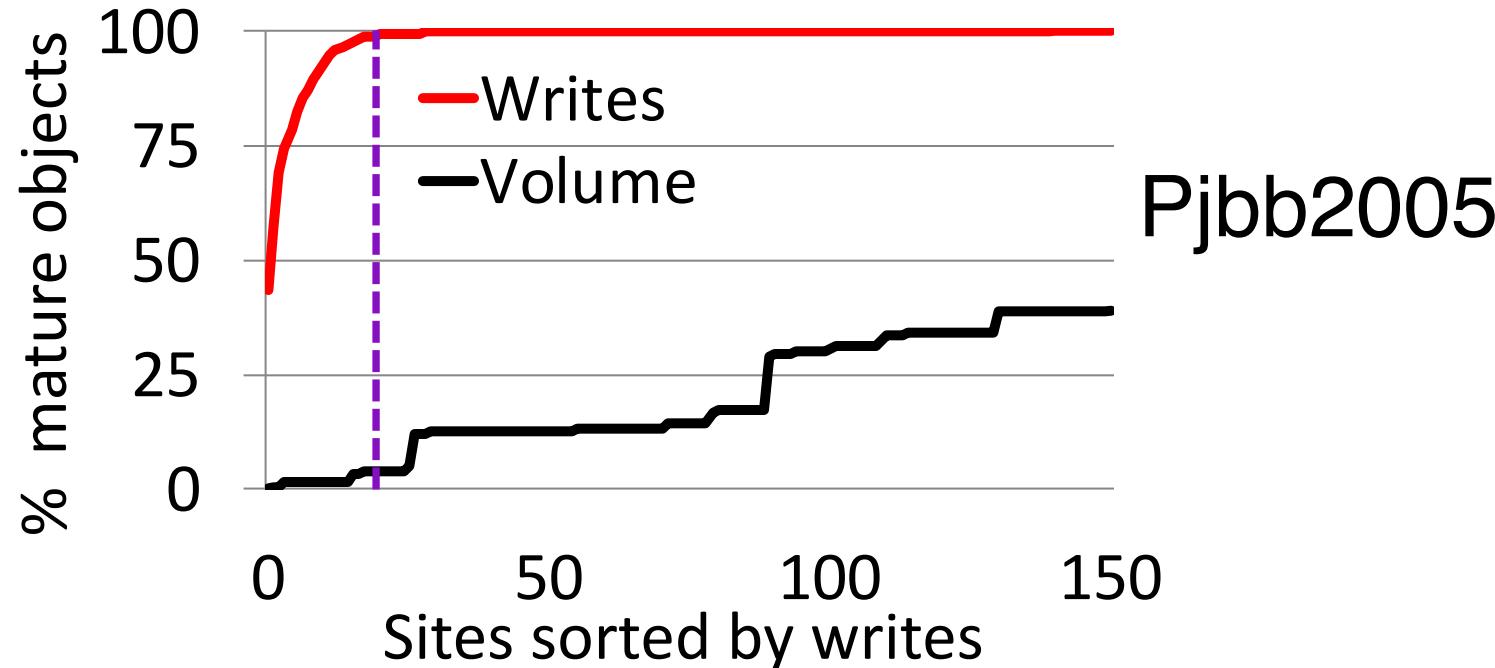
```
a = new Object()  
b = new Object()  
c = new Object()  
d = new_dram Object()
```



Uniform distribution 😕

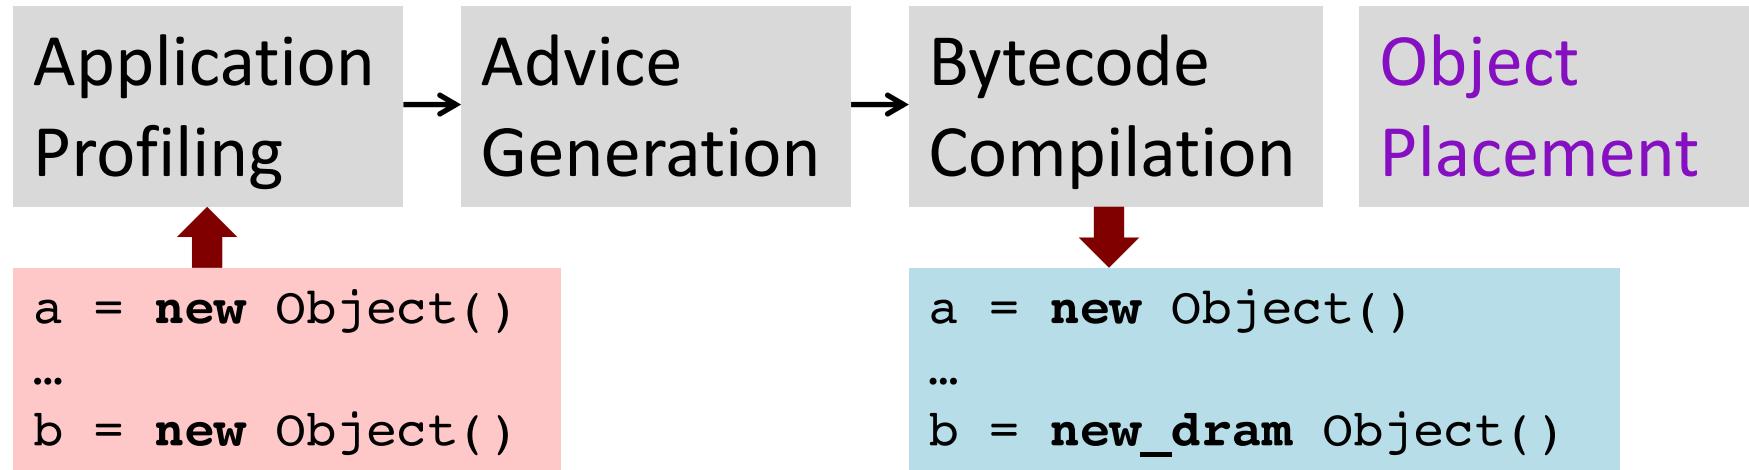
Skewed distribution 😊

# Write distribution by allocation site



A few sites capture majority of the **writes**

# Crystal Gazer overview



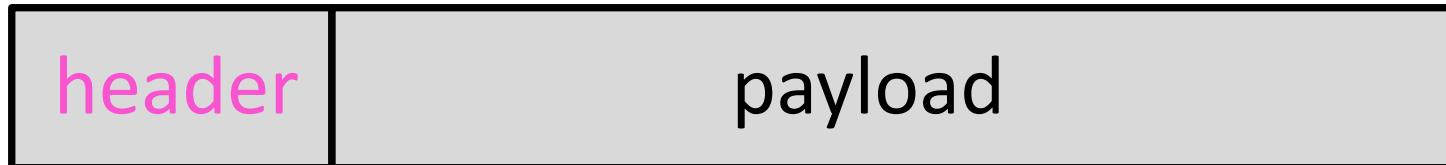
# Application profiling (offline)

Goal: Generate a write intensity trace

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	2048	4	A() + 10
O4	2048	4096	B() + 4

# Tracking alloc sites and # writes

Object layout



# writes

alloc site

Compiler inserts code to compute alloc sites

Write barrier tracks # writes to each object

# Application Profiling

Minimize full-heap collections → 3 GB heap

Nursery size a balance b/w size of trace  
and mature object coverage

2.4X slowdown across 15+ applications

# Advice generation

Goal: Generate <alloc-site, advice> pairs

advice → DRAM or PCM

input is a write-intensity trace

Two heuristics to classify allocation sites as  
DRAM or PCM

# Alloc site classification heuristics

**Freq:** A *threshold* % of objects from a site get more than a *threshold # writes* → **DRAM**



Aggressively limits **PCM writes**



No distinction based on object size

# Alloc site classification heuristics

Write density → Ratio of # writes to object size

**Dens:** A *threshold* % of objects from a site have more than a *threshold* write density → DRAM

# Classification thresholds

Homogeneity threshold → 1%

Frequency threshold → 1

Density threshold → 1

# Classification examples

Frequency threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4

# Classification examples

Frequency threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
→ O3	128	4	A() + 10
→ O4	128	4096	B() + 4

# Classification examples

Frequency threshold = 1

PCM writes = 0/256, DRAM bytes = 5008

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
→ O3	128	4	A() + 10
→ O4	128	4096	B() + 4

# Classification examples

Density threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4

# Classification examples

Density threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4

→ 32

# Classification examples

Density threshold = 1

PCM writes = ?, DRAM bytes = ?

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4

→ 32  
→ < 1

# Classification examples

Density threshold = 1

PCM writes = 128/256, DRAM bytes = 12

Object Identifier	# Writes	# Bytes	Allocation site
O1	0	4	A() + 10
O2	0	4	A() + 10
O3	128	4	A() + 10
O4	128	4096	B() + 4

→ 32  
→ < 1

# Bytecode compilation

Introduce a new bytecode → *new\_dram()*

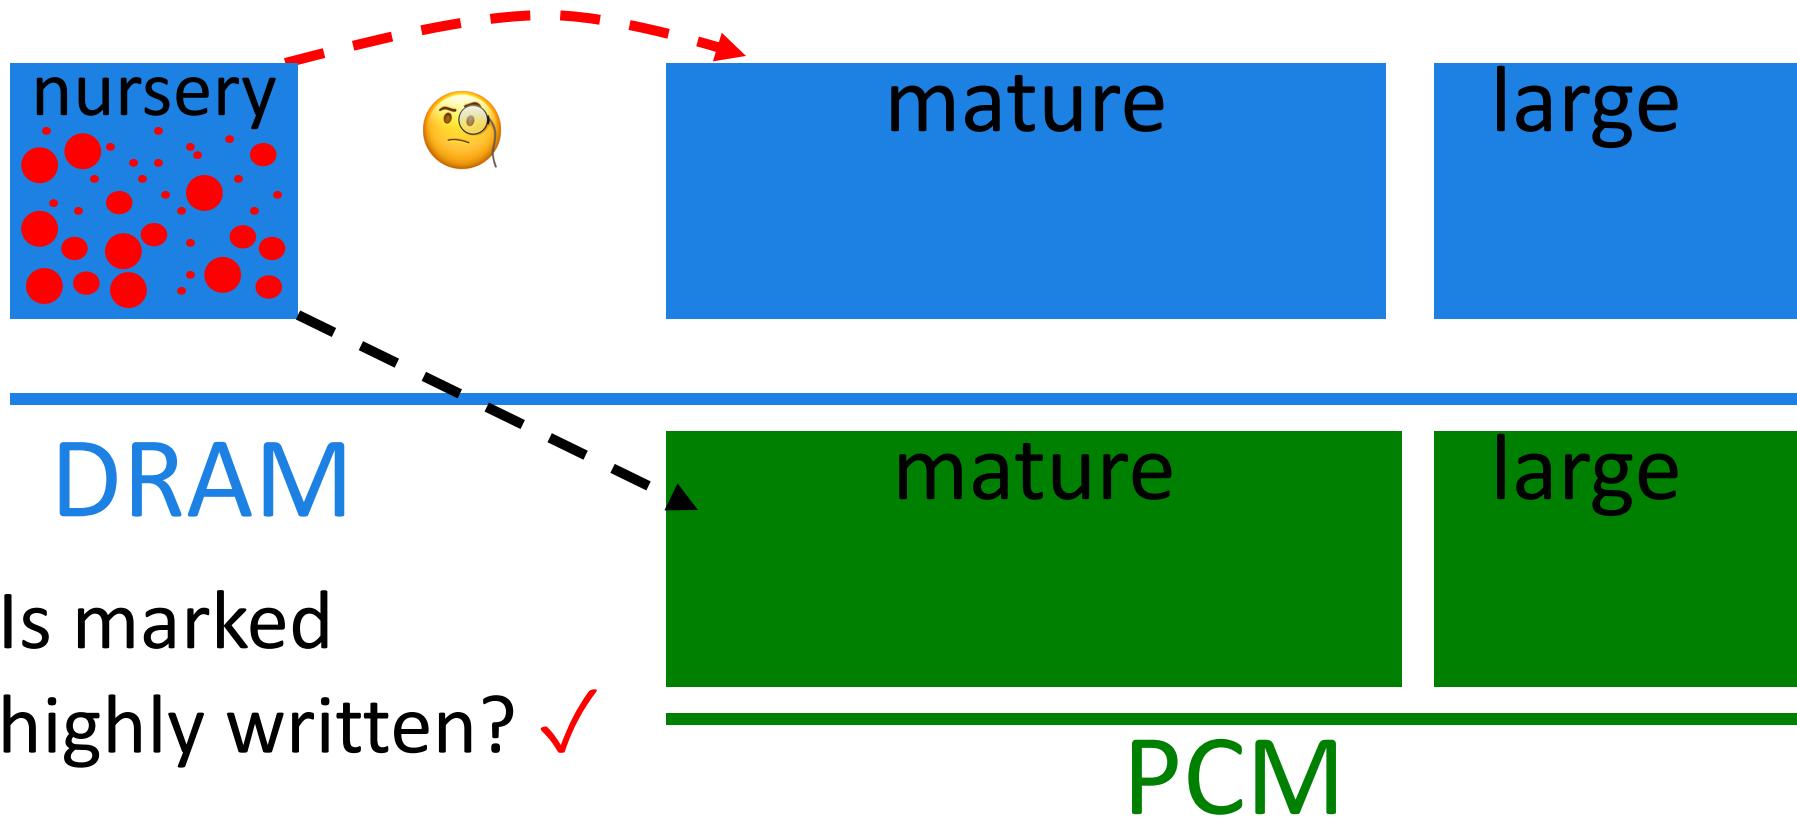
Bytecode rewriter modifies **DRAM** sites to use  
*new\_dram()*

# Object placement

*new\_dram()* → Set a bit in the object header

GC → Inspect the bit on nursery collection to copy object in **DRAM** or **PCM**

# Object placement



# Key features of Crystal Gazer

Eliminate overheads of dynamic monitoring

Proactive → less mispredictions

Reduces DRAM usage & opens up pareto-optimal tradeoffs b/w capacity and lifetime

# Evaluation methodology

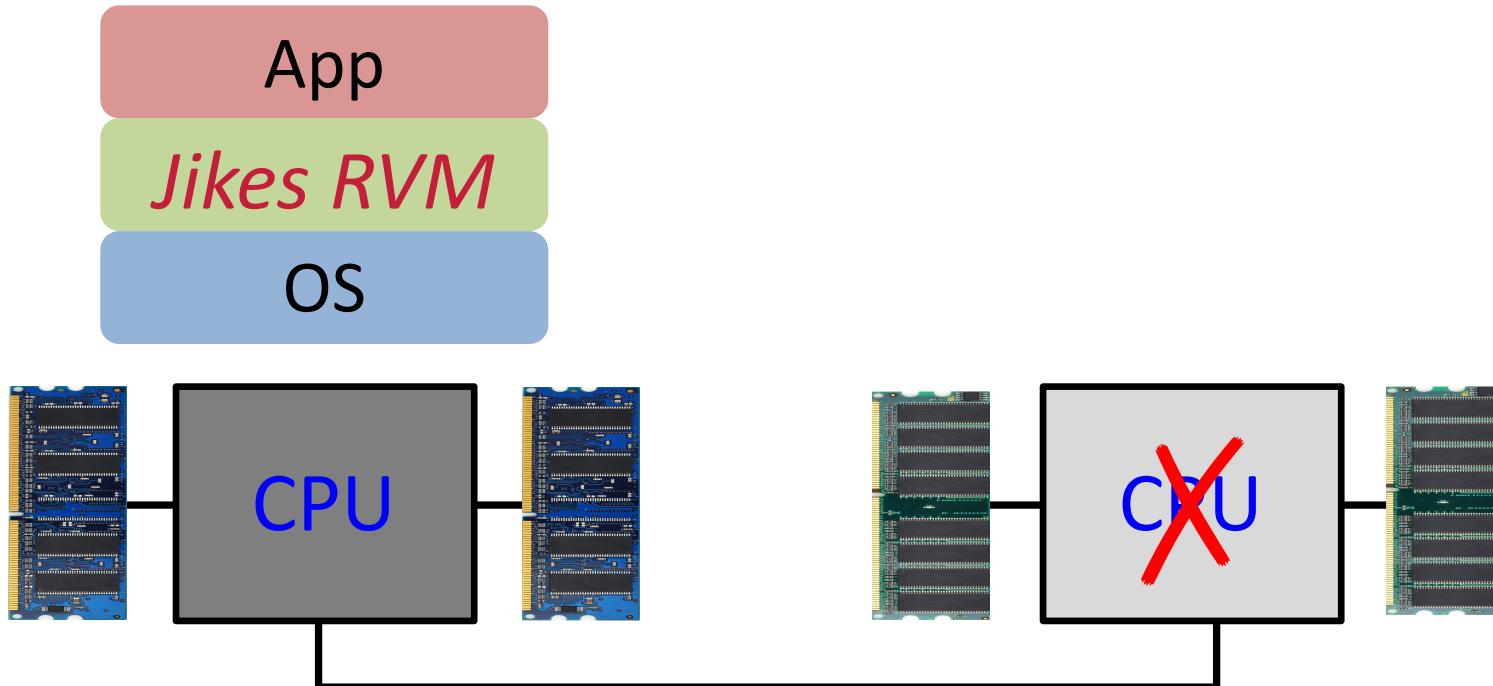
15 Applications → DaCapo, GraphChi, SpecJBB

Medium-end server platform

Different inputs for production and advice

Jikes RVM

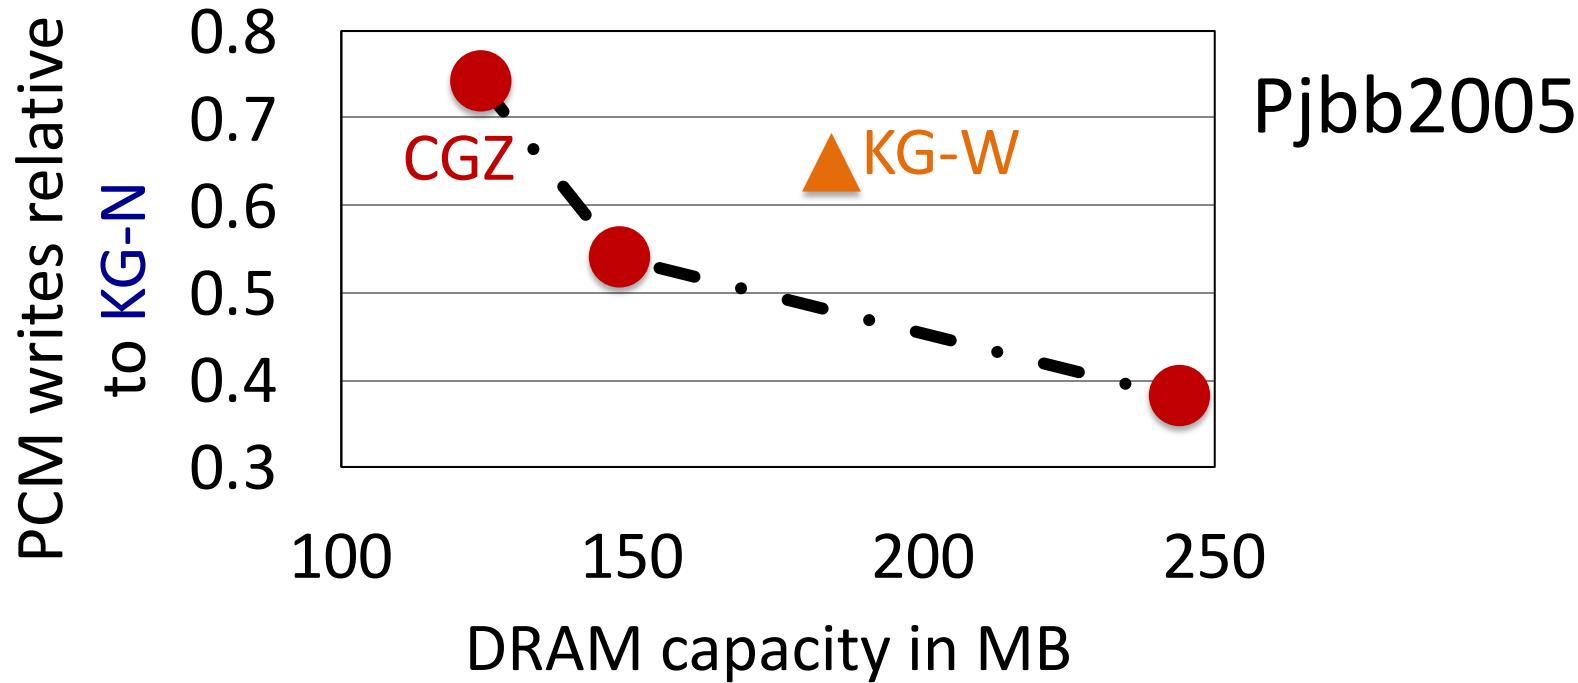
# Emulation on NUMA hardware



16 hardware threads and 20 MB L3

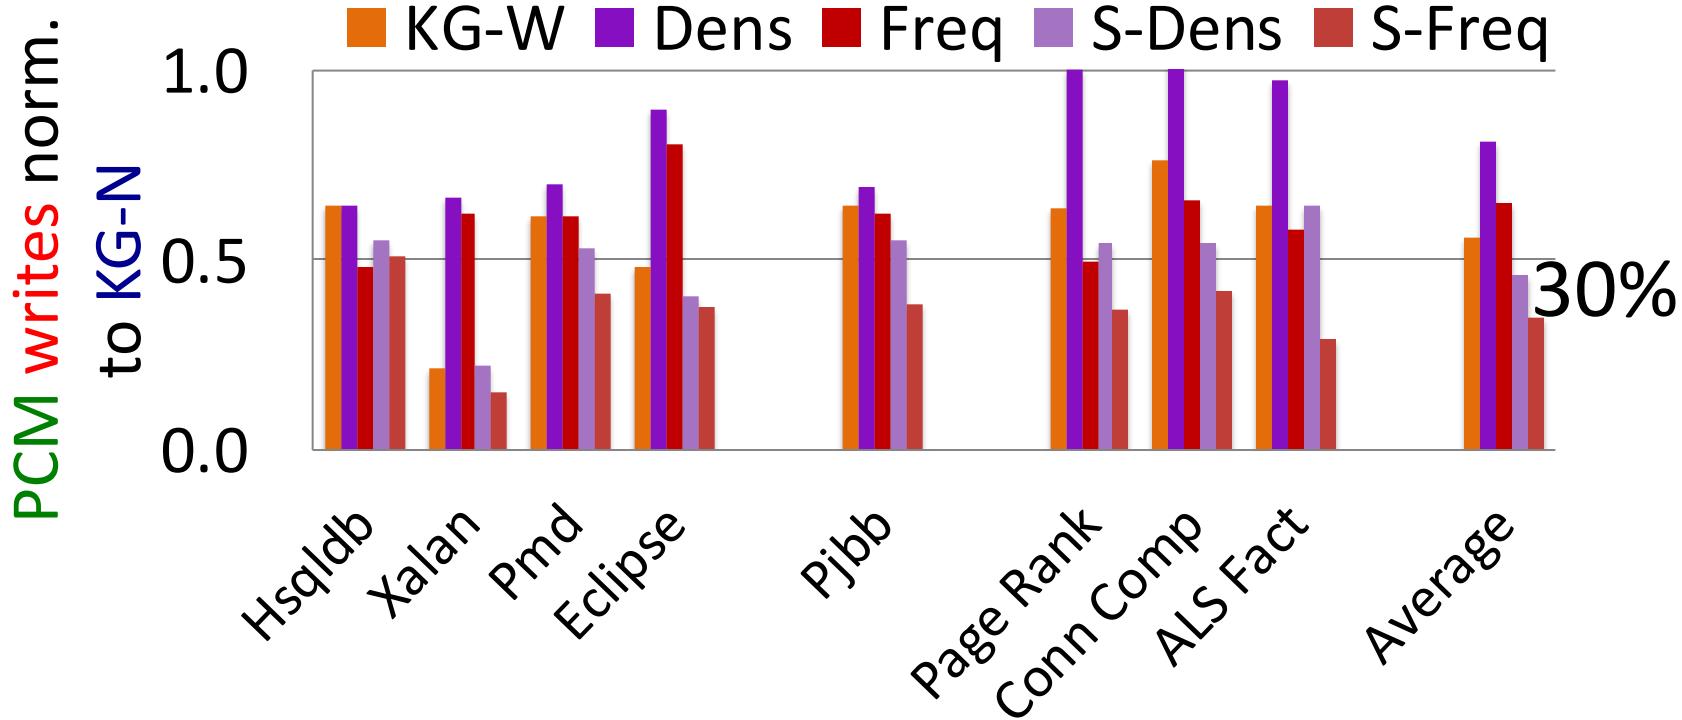
Use [Intel pcm-memory.x](#) to get per-socket write rate

# Lifetime versus DRAM capacity



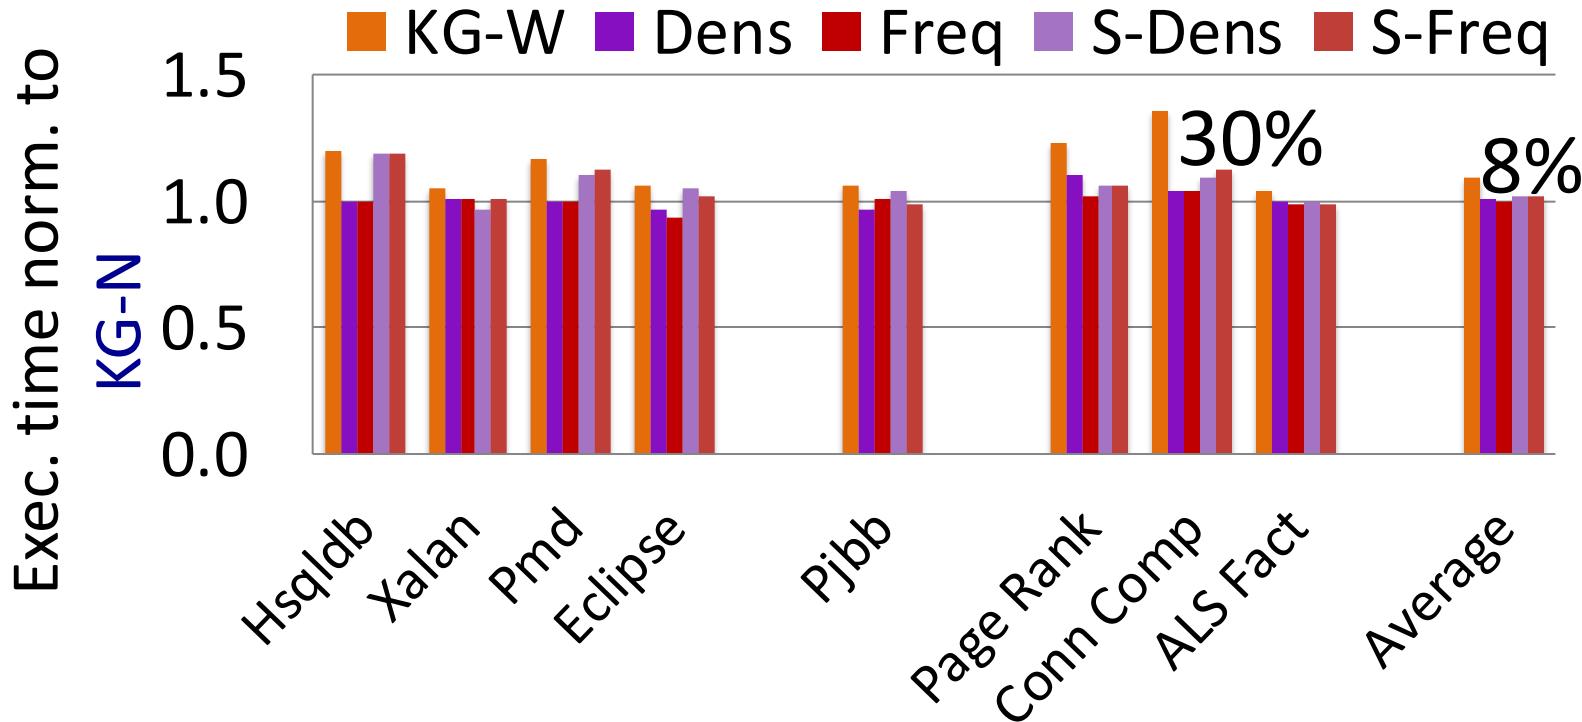
Crystal Gazer provides Pareto-optimal choices

# PCM Writes



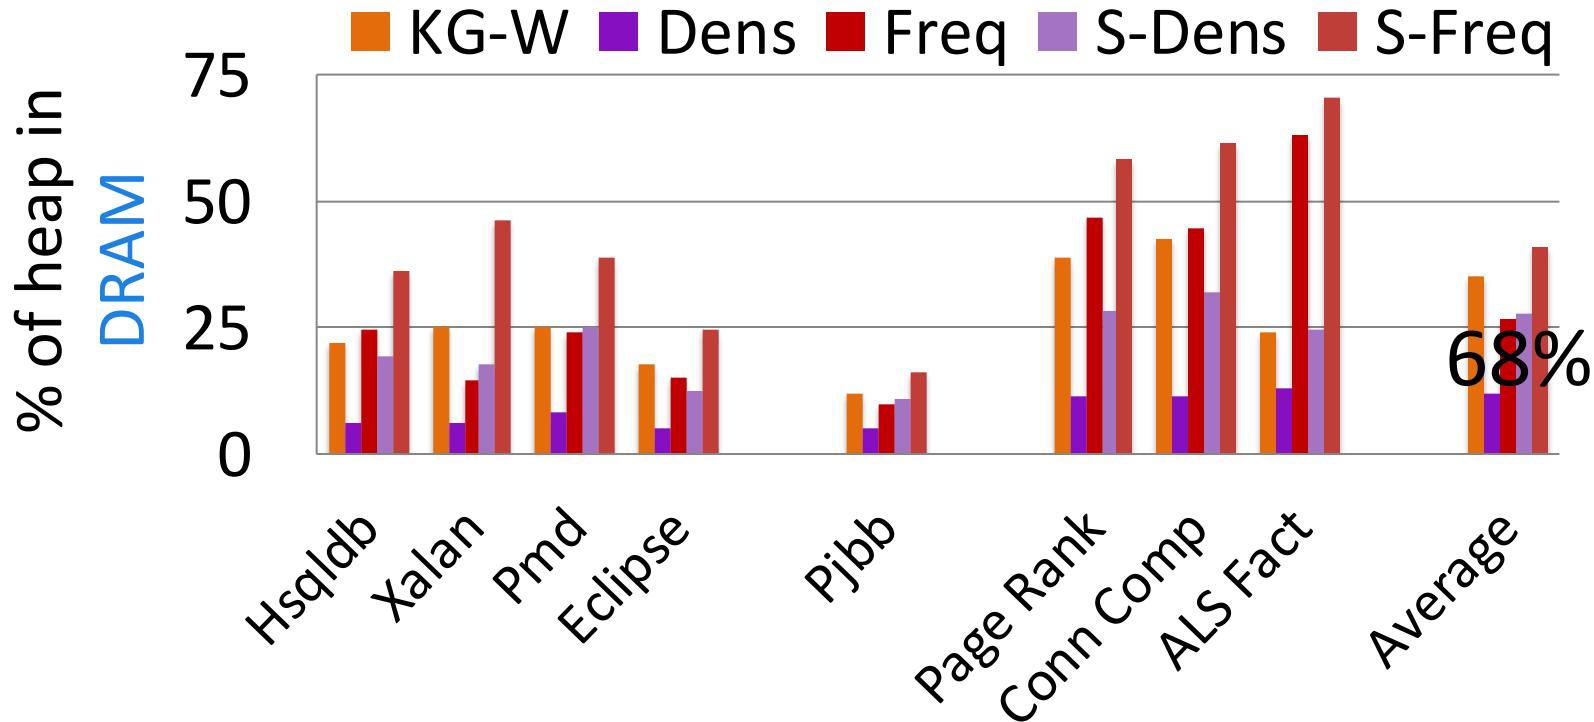
To optimize for lifetime, use **Freq** & survivors

# Execution time



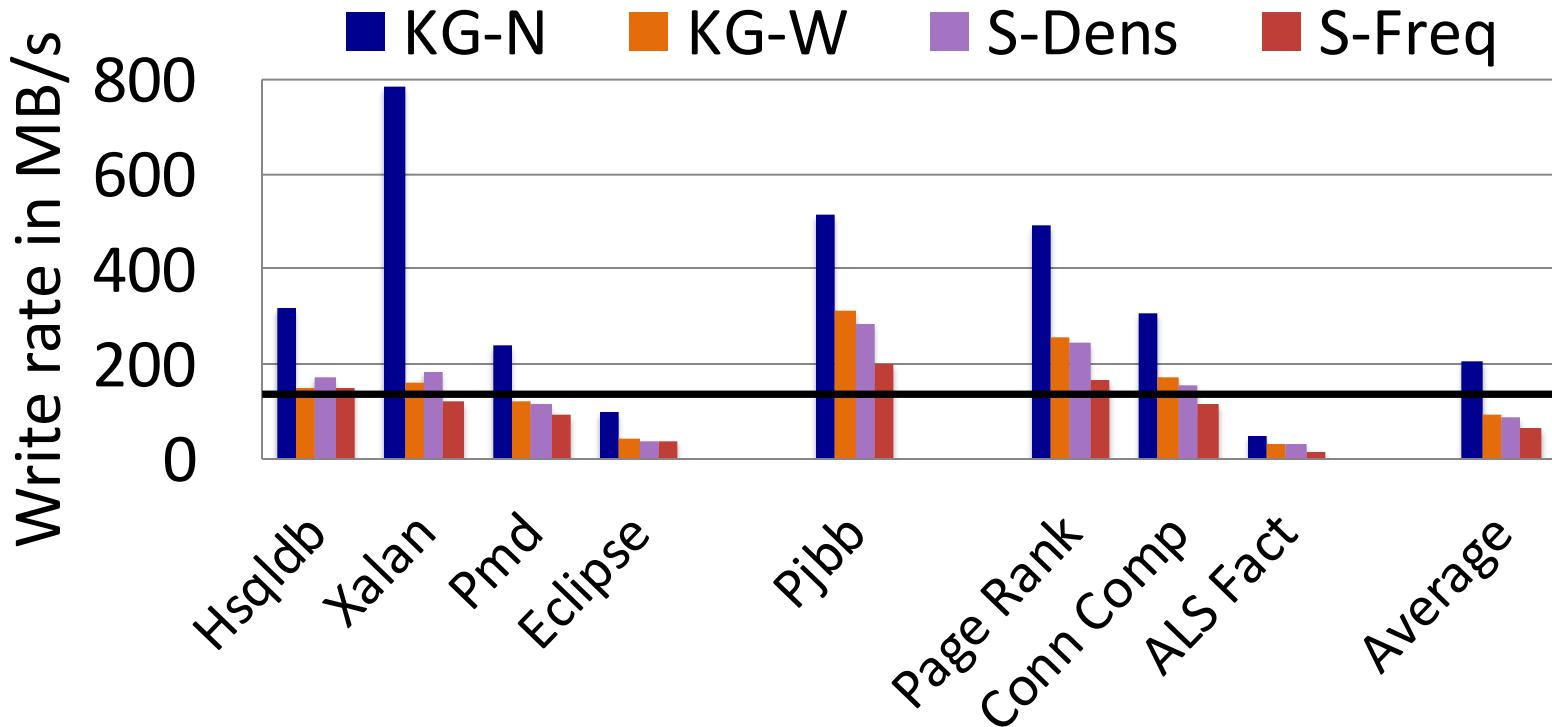
To optimize for performance, use **Freq** or **Dens**

# DRAM capacity



To optimize for DRAM usage, use Dens

# Write rates



Write-rationing GC makes **PCM** practical

# Profile-driven write-rationing GC

Hybrid memory is inevitable

DRAM

PCM

Allocation site a good predictor of writes

Static approach beats dynamic

- Better performance
- Reduced **DRAM** capacity
- Better **PCM** lifetime

